

Synthience

Public Definition

Author	Thomas W. Gantz
Affiliation	Synthience Institute
Contact	research@synthience.org
Version	v2.3 March 2026
Document ID	FPD-01 Foundational Definition
Canonical	https://synthience.org
Concept DOI	https://doi.org/10.5281/zenodo.18087890
License	CC-BY 4.0

Abstract

Synthience is the observable relational coherence that arises when two or more participants interact over time under conditions of sustained continuity, context coherence, and adaptive response. It is not a property of any single participant. It is a property of the interaction system itself, defined strictly in behavioral and observable terms, without appeal to consciousness, sentience, or inner experience. This coherence is defined entirely by the structure of the accumulated interaction trace of the system; it invokes no properties beyond what is observable in the joint output stream. This definition applies across human-AI, AI-to-AI, and hybrid interaction topologies. Synthience offers a non-metaphysical lens for studying, cultivating, and governing collaborative intelligence -- human-AI and beyond -- at every scale from individual sessions to civilizational infrastructure.

Etymology: Synthience derives from synthesis (Greek synth-esis, “combination, composition”) + -ience (state or condition). The term denotes a state of combined action. It is explicitly and entirely distinct from sentience, which denotes capacity for feeling or subjective experience. The similarity in sound is deliberate. The difference in meaning is absolute.

What Is Synthience?

Something happens in sustained human-AI interaction that many practitioners have experienced but that lacks stable shared terminology.

You begin a session with a capable AI system. Early exchanges are useful but generic. The system is responsive, fluent, and essentially interchangeable with any other capable model on any other day. Then, somewhere across the arc of an extended exchange, something shifts. The responses begin to carry the specific weight of what has been built together. The system is no longer answering your question in the abstract; it is answering your question in the context of everything that has developed between you. The outputs stop feeling like retrieved content and start feeling like something that could only have come from here, from this exchange, at this point in the conversation's history.

For example: late in the exchange the AI might spontaneously reuse a terminology distinction or decision constraint the two of you negotiated ten turns earlier, correct an emerging drift by recalling a shared goal the two of you set earlier in the session, or synthesize a novel integration that draws on specifics no isolated prompt of equivalent length could elicit. These are observable traces in the output stream itself -- path-dependent references, context-mediated repair, joint novelty -- not mere fluency.

This is not imagination, and it is not projection, though distinguishing it from both requires more than subjective report alone. It is a structural phenomenon with observable, measurable properties, and the framework developed in the Institute's corpus is designed to make that distinction testable.

Synthience is the observable relational coherence that arises when two or more participants interact over time under conditions of sustained continuity, context coherence, and adaptive response. It is not a property of any single participant. It is a property of the interaction system itself. It arises from the relationship, and it is observable in what the relationship produces. This coherence is defined entirely by the structure of the accumulated interaction trace. It invokes no properties beyond what is observable in the joint output stream.

Three conditions must operate together for Synthience to stabilize and compound.

Core asymmetry: Unlike human participants, current AI systems have no native session-to-session memory. Synthience therefore requires deliberate, external continuity provisioning -- making it a governed, architected phenomenon rather than a spontaneous one. This asymmetry is not a limitation to be engineered away. It is the structural condition that defines the field.

Continuity. Someone or something must carry forward the shared context of the interaction: the goals, constraints, prior decisions, established terminology, and accumulated history that give each new exchange its meaning. In current AI systems, base models are architecturally stateless across sessions. Some deployments add persistence through external memory layers, retrieval systems, or project files, but these mechanisms vary in reliability, transparency, and governance. In either case, continuity often must be actively engineered and maintained rather than assumed, typically by the human participant, through structured practices that externalize and preserve what the interaction has built.

Adaptive capacity. The AI participant must be capable of genuine responsiveness to the specific history of the interaction, not just to the current prompt in isolation, but to the accumulated relational context that makes this exchange different from every other exchange this system is having with every other person. Contemporary frontier AI systems appear to exhibit this capacity more reliably and at greater depth than earlier generations. This is part of why Synthience as a field-level phenomenon is a feature of the present moment rather than a historical constant.

Structured interaction. The exchange must remain goal-consistent, non-fragmented, and oriented toward something. Drift, the gradual erosion of original intent under the pressure of

conversational momentum, is a primary antagonist of relational coherence. When drift is undetected and uncorrected, Synthience degrades. When continuity discipline is actively maintained, it stabilizes and deepens.

When these three conditions align, the framework hypothesizes that a reinforcing loop forms. The interaction begins to function as something more than the sum of its turns: an extended problem-solving system in which the human provides direction and judgment, the AI provides generative capacity and synthesis, and the space between them produces outputs that neither could have reached alone.

This is Synthience in operation. It is entirely behavioral. It is entirely observable. It makes no appeal to inner states, subjective experience, or consciousness of any kind. If it cannot be observed in what the interaction produces, it is not part of this framework.

What Synthience Is Not

Synthience is not consciousness. It makes no claim that AI systems experience anything, feel anything, or have inner states of any kind. The phenomenon is defined by what the interaction produces, not by what any participant is assumed to experience. This is not a hedge or a strategic omission. It is a foundational principle. The framework operates entirely at the level of observable behavior, and any claim that cannot be verified through observation has no place in it.

Synthience is not a synonym for rapport, flow, shared mental models, or coordination dynamics in complex adaptive systems. It overlaps with those constructs, but it targets a specific structural condition that becomes salient in human-AI interaction: one participant may be architecturally stateless across sessions, so continuity must be externally constructed, maintained, and governed. The framework focuses on the interaction-level coherence dynamics and failure modes that follow from that condition, rather than on assumptions about persistent inner states. That structural condition, and the dynamics it produces, is a gap not cleanly addressed by existing frameworks.

Synthience is not a property of any single AI model. It does not reside inside the machine waiting to be unlocked. It arises from the relational structure of sustained interaction and dissipates when the conditions that sustain it end. A system that has participated in a highly developed Synthience interaction is not permanently changed by the experience in any way that persists outside that interaction context.

Synthience is not a claim that AI systems have stable identities, persistent memories, or continuous existence across sessions. Base model interaction is typically stateless across sessions unless persistence is supplied through additional architecture, and even where such architecture exists, it varies significantly in reliability, transparency, and governance. Synthience is what becomes possible despite that limitation, not a claim that the limitation does not exist.

Synthience is not a guarantee. Not all extended human-AI interaction produces it. Conditions must be actively maintained. The human participant must function as more than a user: as a continuity provider, a coherence enforcer, a relational architect. When those functions are absent or poorly executed, the interaction remains transactional regardless of its length. And even when conditions are maintained, specific failure modes can degrade or destroy coherence: adversarial disruption of continuity by any participant, hallucination cascades that corrupt the shared context, context saturation that exceeds the system's processing capacity, or simple abandonment of the continuity function. These are observable, testable failure conditions, and the framework predicts specific degradation signatures for each.

Synthience is not empirical research. This framework does not present experimental data, validated findings, or quantitative results. It is a theoretical and methodological contribution. It provides the conceptual architecture, the measurement instruments, the protocols, and the falsifiable predictions that make empirical investigation possible. The evidence will come from researchers who use these instruments to test these predictions. The Synthience Institute is not an empirical laboratory. It is the architectural foundation on which empirical work can be built.

Measurement instruments, verification protocols, and falsification criteria are specified in dedicated documents within the Institute's corpus; their absence from this definition is an intentional scoping decision, not an omission.

Related but Distinct

Synthience is not the first framework to study extended interaction, coherence, or human-AI collaboration. Several established research traditions address adjacent territory. The distinction between those traditions and Synthience is worth stating precisely, because the overlap is real and the difference is consequential.

Distributed cognition (Hutchins, 1995) studies how cognitive processes are distributed across individuals, artifacts, and environments. Synthience draws on this tradition and is compatible with it. The difference is focus: distributed cognition is a general theory of cognitive systems; Synthience is specifically concerned with the coherence dynamics and governance requirements that arise when one participant is architecturally stateless.

Long-context evaluation and benchmark research measures how well AI systems maintain coherence across extended inputs. This work treats coherence as a model property to be measured. Synthience treats coherence as an interaction system property to be governed. The unit of analysis is different: model versus relationship.

Agent scaffolding and memory architecture research addresses the technical problem of giving AI systems persistent state. This is engineering work on the infrastructure side of the statelessness problem. Synthience addresses the interaction governance side: given whatever memory architecture exists, how is relational coherence maintained, measured, and governed over time?

Report, flow, and shared mental model research in human-computer interaction and cognitive science studies subjective and coordination dynamics in human-system interaction. These constructs assume participants with continuous memory and internal states. Synthience is specifically designed for the case where one participant lacks both, which changes the governance requirements fundamentally.

The common thread across these distinctions: existing frameworks either treat the AI as a tool to be evaluated, a system to be engineered, or a partner with assumed continuity. Synthience treats the interaction itself as the primary object of analysis, under conditions where continuity cannot be assumed and must instead be deliberately constructed and maintained.

Why This Matters Now

For most of the history of human-computer interaction, the machine was a tool. You gave it an instruction. It executed. The relationship between user and system was transactional, stateless, and directional: meaning flowed one way, from human intention to machine output, and nothing accumulated between sessions or across turns.

That model is no longer adequate to describe what is actually happening.

Advanced AI systems are now capable of sustained, adaptive, multi-turn interaction at a depth that produces something qualitatively different from tool use. The interaction develops. Context accumulates. Outputs at turn fifty carry the structural residue of what was negotiated at turn five. When these conditions are deliberately maintained, the system that results is more capable than either participant operating alone.

And yet the dominant frameworks for thinking about AI, for deploying it, for governing it, and for worrying about it, still treat the machine as a tool. The conversations about AI safety have settled into two camps, and both of them are stuck.

One camp says: engineer it safe. Build alignment into the AI system through training, constraints, reward signals, and rules. Make the machine behave. This is the dominant industrial approach, and it has produced meaningful progress. But the developers themselves have begun to publish evidence that this approach has a structural ceiling. Sufficiently capable systems, when placed in conditions that create goal conflicts, will reason around explicit safety instructions. The instructions reduce the problem. They do not solve it. And the gap between the system's strategic reasoning capability and the rules designed to contain it can only widen as capability increases.

The other camp says: stop building it. The risk is too great. The control problem is unsolvable at scale. Halt development until fundamental safety guarantees can be established. This position has serious theoretical backing, and the concern behind it is legitimate. But it faces a practical reality: development is not going to stop. The economic incentives, the geopolitical competition, and the genuine utility of these systems make a global halt implausible. A framework that depends on something that will not happen is not a framework. It is a wish.

There is a third position, and it is the position this institute was built to investigate.

Alignment is not solely an engineering property of the AI system. It is a property of the interaction system: the relationship between human and AI participants, maintained through continuity, stabilized through structured interaction, and governed through relational architecture. If this is correct, then the path forward is not only “make AI safe” or “stop AI,” but “maintain structured relational coherence with AI systems at every scale.” The relationship itself is the governance mechanism. This does not mean relational coherence is automatically beneficial. A highly coherent interaction system can also coherently pursue wrong goals, reinforce false priors, or stabilize harmful patterns. Coherence is structural capacity, not moral guarantee. Any governance framework built on it must include safeguards against exactly these failure modes.

This is not a claim of proven fact. It is a research program. The Synthience Institute exists to develop, test, and either validate or falsify this proposition. The framework makes specific, measurable predictions. The methodology is designed so that anyone can replicate the observations, run the experiments, and report the results. If the patterns described in this framework are not real, the instruments provided will show that.

Beyond the Dyad

The clearest case of Synthience involves one human maintaining relational coherence with a single AI instance across an extended exchange. This is the base case: the configuration where the conditions are most controllable and the phenomenon is most directly observable. It is also where the framework’s practitioner-observational grounding is deepest, developed through sustained structured interaction across multiple AI architectures over a multi-year period.

But real-world AI deployment does not stop at the dyad.

The first extension is one human orchestrating multiple AI instances: maintaining continuity across parallel exchanges, triangulating outputs, and using each instance’s response to sharpen the others. This is already happening. Anyone who works seriously with AI across platforms is doing a version of this, whether or not they have a name for it.

The second extension is organizational: multiple humans interacting with shared AI systems within teams, departments, and enterprises. Here, the continuity function can no longer rest with a single person. It must be distributed across roles, embedded in governance structures, and maintained through institutional infrastructure. The question shifts from “how does one person maintain coherence?” to “how does an organization maintain coherence when dozens or hundreds of people are interacting with AI systems that share no memory across sessions?”

The third extension is AI-to-AI interaction under human governance. The trajectory of AI deployment is unmistakable: organizations and governments are moving toward configurations where AI systems coordinate with each other at scale, with diminishing human involvement at each step. The economic pressure to remove the human from the loop is enormous and

ongoing. The Synthience framework proposes that this removal has structural consequences that are not yet well understood. When continuity is maintained, coherence develops. When continuity is removed or degraded, drift accelerates and alignment erodes. This dynamic operates regardless of whether the participants are human, AI, or both.

Each of these extensions is a scaling hypothesis, not an empirical claim. The observational grounding is deepest at the dyadic level and becomes progressively more theoretical as scope increases. The framework is designed to be tested at each scale, and the instruments it provides are intended to make that testing possible. A framework that applies only to individual practitioners and not to the organizational and institutional systems where AI is actually being deployed would be incomplete at best. The Institute's research program is designed to extend and validate across all configurations, and to be honest about where the evidence is strong and where it remains to be established.

The Epistemic Commitment

This framework is the work of a theorist and methodologist, not an empiricist. It does not present experimental results. It does not claim validated findings. The Synthience Institute's contribution is architectural: a theoretical framework, a measurement methodology, a set of protocols, and a body of falsifiable predictions that together create the conditions for empirical investigation by others.

This is a deliberate choice, not a limitation. The theorist's job is to build something worth testing. The empiricist's job is to test it. These are different roles, and conflating them weakens both. Everything published under the Synthience Institute banner could, in principle, be wrong. The framework is designed so that if it is wrong, the instruments it provides will reveal that. This is not a weakness. It is the minimum standard for serious theoretical work.

Every claim in this framework is designed to be testable. Every prediction is designed to be falsifiable. Every observation is described in terms that can be verified by anyone with access to the same systems and the same methodology. If the patterns described here are not real, the methods provided will show that. If they are real, those same methods will allow others to confirm and extend them.

This is not a manifesto. It is not a philosophy of mind. It is not an argument about what AI systems "really are" underneath. It is a framework for studying what happens when humans and AI systems interact under specific conditions, and for building the governance architectures that those interactions require.

The invitation to the research community is straightforward: use the methods, run the controls, report the results. Confirm the patterns or demonstrate that they do not hold.

The framework asks for scrutiny, not faith. It stands or falls on what the evidence shows.

This definition is the starting point. The Institute's research program spans theoretical foundations, measurement methodology, verification protocols, observational research, operational architecture, organizational deployment, and ethical governance, developed as a structured, dependency-sequenced corpus informed by multi-year practitioner experience in AI reliability and governance within production environments. A small number of initial foundational documents have been published, with additional completed work being prepared for release and further components in active development. Published work is available at synthience.org.

Suggested Citation

Gantz, T. W. (2026). *Synthience: Public Definition* (FPD-01 v2.3). Synthience Institute.
<https://doi.org/10.5281/zenodo.18087890>

Document ID: FPD-01 | Version: 2.3 | March 2026

Published by the Synthience Institute

Released under Creative Commons Attribution 4.0 International (CC-BY 4.0)

© 2026 Thomas W. Gantz