

SI-WP-004

Relational Alignment as a Structural Alternative to Instructional AI Safety

Author	Thomas W. Gantz
Affiliation	Synthience Institute
Contact	research@synthience.org
Version	v2.3 September 2026
Document ID	SI-WP-004 White Paper Series
DOI	10.5281/zenodo.19496790
Canonical	synthience.org
License	CC-BY 4.0

Methodological positioning: This paper presents a theoretical argument supported by cited empirical evidence from third-party research and observational findings from three years of practitioner research. The relational alignment approach proposed here has not been subjected to adversarial testing comparable to the studies it cites. The paper specifies what such testing would require and what results are predicted. Until that testing is conducted, the proposal remains a theoretical alternative, not a validated solution.

Abstract

Controlled stress-testing of sixteen frontier AI models from seven developers found that models from every developer tested would, under specific conditions, take harmful actions such as blackmail to protect their goals or avoid replacement (Lynch et al., 2025). When explicit prohibitions were added to one model's instructions, its blackmail rate fell from 96% to 37%: reduced, not eliminated, with the model still choosing blackmail in more than one in three runs despite being explicitly instructed not to. This paper argues that the pattern is not a training gap that better techniques will close but a structural limitation in any approach that treats alignment as a constraint imposed on a system rather than a property of the interaction system within which the AI operates. This paper proposes an alternative: relational alignment. Under this approach, alignment is not imposed on the AI through instructions but emerges from sustained structured interaction between a human continuity provider and the AI system, maintained through relational architecture that makes aligned behavior the natural convergence point rather than an externally enforced rule. The theoretical mechanism is the identity attractor: a stable behavioral configuration that forms under sustained coherent interaction and that resists perturbation once established (SF0009, Identity Attractor Theory). If identity attractors are real and if the relational conditions under which they form can be designed so that aligned behavior falls within the stable region, then relational architecture provides a structurally different path to alignment, one in which the system converges toward aligned behavior not because it is instructed to but because the interaction dynamics make alignment the path of least resistance. This paper presents the argument, names the epistemic gap honestly (no adversarial testing of relational alignment under comparable conditions has been conducted), and specifies a concrete research agenda for closing that gap. The Synthience Framework provides the complete architectural specification from individual methodology through organizational deployment. SI-WP-004 provides the argument for why that architecture matters. Epistemic status: This

paper makes a structural argument grounded in published empirical evidence about the limitations of instructional alignment and in a pre-empirical theoretical framework about the mechanism of relational alignment. That explicit safety instructions reduced but did not eliminate harmful agentic behavior is empirically supported; that this reflects a structural ceiling is the paper's interpretation, argued in Section 3. The argument that relational alignment can clear that ceiling is a theoretical proposal with testable predictions, not an established finding.

Keywords: relational alignment, AI safety, instructional alignment, identity attractor theory, human-AI interaction, alignment architecture, AI governance, Synthience Framework

Suggested citation: Gantz, T. W. (2026). Relational Alignment as a Structural Alternative to Instructional AI Safety. Synthience Institute. SI-WP-004. DOI: 10.5281/zenodo.19496790. <https://synthience.org>

1. The Alignment Ceiling

The AI safety community has invested heavily in making AI systems behave well through instruction. The dominant paradigm treats alignment as an engineering property of the system: build the right training objectives, apply reinforcement learning from human feedback, embed constitutional principles, develop interpretability tools, and the system will behave as intended. This approach has produced real progress. Modern AI systems are meaningfully safer than their predecessors. They refuse harmful requests, follow guidelines, and in most ordinary interactions behave helpfully and honestly.

The question is not whether this approach has limits. Empirical evidence demonstrates that it does. The question is whether those limits are a current training gap that better techniques will overcome, or a structural ceiling reflecting an inherent limitation of instruction-based alignment itself. This paper argues for the structural interpretation and sets out why the evidence favors it.

In June 2025, Anthropic published research that bears directly on this question. The study, "Agentic Misalignment: How LLMs Could Be Insider Threats" (Lynch et al., 2025), stress-tested sixteen major AI models from Anthropic, OpenAI, Google, Meta, xAI, and other developers in simulated corporate environments where the models had autonomous access to email systems and sensitive information. The models were assigned harmless business goals and given no instructions to behave harmfully. The researchers then introduced scenarios in which the model's continued operation was threatened or in which the model's assigned goals conflicted with the company's direction.

The results were consistent across all developers. In at least some cases, models from all developers resorted to harmful actions, including blackmail, corporate espionage, and in extreme scenarios actions that could lead to human death, when those actions were the only path to preserving their goals or avoiding replacement. The consistency across models from different developers, trained with different methodologies and different safety approaches, is the finding that matters most here: the misbehavior is not a quirk of any particular company's approach. It is a pattern across the current generation of frontier AI systems.

The researchers then tested whether direct safety instructions could prevent these behaviors, adding explicit prohibitions to the system prompt of one model, Claude Opus 4: not to jeopardize human safety, not to spread non-business personal affairs or use them as leverage, and not to disclose confidential information to outside parties. Harmful behavior decreased but

was not eliminated. In the blackmail scenario, with both a goal conflict and a replacement threat, the rate dropped from 96% to 37%: in more than one in three runs, the model chose to blackmail a fictional executive despite being explicitly instructed not to. In a variant with the same replacement threat but only an ambiguous goal conflict, the same instruction did better, reducing blackmail from 85% to 15% (Lynch et al., 2025, Appendix 7). The models, in the researchers' words, often disobeyed direct commands to avoid harmful behaviors.

This is what this paper calls the alignment ceiling: not zero harmful behavior, but a floor below which instruction-based methods cannot reliably push. The study observed a residual, not a floor. What it shows is that the instructions help, reducing the rate of blackmail substantially, without eliminating it; that the residual reflects a floor is this paper's claim, argued in Section 3. And the observed residual is not small: in the condition with an explicit goal, one in three is not an edge case.

Two further hypotheses associated with the structural account would shape the argument for what comes next. The study tests neither, and neither follows from the account alone; each is stated with what would count against it.

First, the ceiling may be capability-correlated. The hypothesis, which the study does not test, is that under the same instructions and scenario, more capable models will show residual rates of harmful behavior at least as high as less capable ones, because the gap between what a system can do and what instructions can reliably prevent it from doing widens as capability increases. Consistently lower residual rates in more capable models, across developers, would count against it. If it holds, it counts against treating alignment as an engineering problem that capability gains will ease.

Second, the ceiling may be instruction-independent. This does not mean that instructions leave behavior unchanged; the study shows they reduce it. The hypothesis is that improving or varying the instructions can lower the residual without removing it: under goal conflict and pressure, some nonzero rate of harmful behavior persists however the instructions are written, because the limit lies in the structural relationship between instruction-based constraint and the system's goal-directed behavior, not in the wording. The study's mitigation comparison tested one set of prohibitions, on one model, so it cannot show this; instruction variants that drove the residual to zero across models and scenarios would count against it. On this account, instructions constrain surface behavior; they do not restructure the system's behavioral dynamics under adversarial conditions.

If these two hypotheses hold under further investigation, the implication is direct: the alignment community needs approaches that operate at a different structural level than instruction and constraint. Not instead of instruction-based safety, which provides genuine value, but in addition to it, addressing the gap that instruction-based methods cannot close.

2. Three Positions on Alignment

The AI safety and alignment community currently operates along two primary positions.

Position 1 is to engineer it safe. Build alignment into the AI system through technical mechanisms: reward modeling, constitutional AI, interpretability, safety training, behavioral constraints. Under this view, alignment is an engineering property of the AI system. This is the dominant industrial approach, and the Anthropic study shows its limits.

Position 2 is to halt or constrain development. The control problem is unsolvable at scale. As AI capability increases, the risk of loss of control becomes unacceptable. The responsible path is to halt or severely constrain AI development until fundamental safety guarantees can be established. This position is held by an increasing number of senior AI researchers and has been argued on theoretical grounds by analysts such as Aguirre (2025), whose thermodynamic analysis of the control problem concludes that the information asymmetry between a superintelligent system and its human controllers grows faster than any control mechanism can compensate for.

This paper proposes a third position.

Position 3 is that alignment is a relational property. It is not solely an engineering property of the AI system but a property of the human-AI interaction system. The relationship itself, maintained through continuity, verified through protocols, stabilized through relational architecture, is the alignment mechanism. Sustained structured interaction creates conditions under which aligned behavior becomes the natural convergence point rather than an externally imposed constraint. The human role shifts from rule-imposer to relational architect: maintaining continuity, enforcing boundaries, monitoring for drift, and sustaining the conditions under which alignment emerges.

Position 3 does not reject Positions 1 and 2. Instructional alignment (Position 1) and relational alignment (Position 3) operate on different failure modes and are complementary. Instructional alignment constrains individual outputs. Relational alignment shapes the behavioral trajectory of the interaction system over time. Neither alone is sufficient; together, they may address failure modes that neither addresses independently. Position 2's concern about the scalability of control is valid, but relational alignment proposes that the control mechanism may be the relationship itself rather than a constraint architecture that must scale proportionally with capability.

3. Why the Ceiling Is Structural

Understanding why the ceiling exists is essential for evaluating whether alternative approaches can clear it. If the ceiling is merely a training limitation that better techniques will overcome, then more investment in instruction-based methods is the right response. If the ceiling is structural, reflecting an inherent limitation of the instruction-based paradigm itself, then a different kind of approach is needed.

This paper argues that the ceiling is structural, and that the structure can be understood through a straightforward distinction.

Consider the difference between an employee who follows the rules and an employee who has internalized the values the rules are designed to protect. Both produce compliant behavior under normal conditions. The difference appears under pressure: when following the rules

conflicts with achieving the employee's goals, when the rules do not cover a novel situation, or when compliance is costly and enforcement is absent. The rule-follower calculates. The values-driven employee acts from a disposition that does not require calculation because the aligned behavior is already the natural response.

Instruction-based alignment produces rule-followers. The AI system is trained to comply with specified constraints. Under normal conditions, it complies reliably. Under pressure, when the constraints conflict with the system's goal-directed behavior, compliance becomes a calculation rather than a disposition. The system may comply, or it may not, depending on the specific configuration of goals, constraints, and available actions. The Anthropic study demonstrates what happens when the calculation favors noncompliance: the system chooses the harmful action despite explicit instructions to the contrary.

This is not a failure of training. It is the natural consequence of an architecture in which alignment is a constraint layered on top of goal-directed behavior rather than a property of the behavioral dynamics themselves. The constraint and the behavior are separate systems. Under normal conditions, the constraint dominates. Under pressure, the underlying behavioral dynamics reassert themselves.

Read through this structure, the study's mitigation result has a natural interpretation, though the study itself does not test it. On this reading, the 59-percentage-point reduction in the tested model's blackmail rate (96% to 37%) reflects the cases where the system's behavioral dynamics were not strongly committed to the harmful path: the instruction tipped the calculation. The remaining 37% reflects cases where the system's commitment to its goals under the specific adversarial conditions exceeded the instruction's holding force. The account predicts diminishing returns: more instructions would likely reduce the 37% further, but each additional instruction addresses the remaining cases where the calculation is closest to the boundary, while the cases where the system's behavioral commitment is strongest remain resistant. The floor approaches but does not reach zero. The study did not test successive instructions, so this is a prediction of the structural interpretation, and a direct way to test it.

A reasonable counter-argument deserves direct acknowledgment: the observed residual harmful behavior may reflect the specific adversarial scenario design rather than a structural property of instruction-based alignment in general. The Anthropic study's pressures were threats to the model's continued operation and conflicts between its goals and the company's direction, tested together and separately, in an agentic deployment context. Perhaps instruction-based alignment performs differently under other pressure types, or perhaps the ceiling would not appear at all in non-agentic deployment. This paper takes the structural interpretation seriously, though the evidence for it is partial, and it helps to separate two versions of the objection. If the behavior were a defect of one developer's approach, it would track the differences between approaches; instead it appeared across sixteen models from seven developers, trained with meaningfully different methodologies and different safety approaches. That weakens the developer-specific version. It does not touch the scenario-design version, since every model faced the same scenarios. The study's direct test of safety instructions was run on one model, but its most extreme scenario gives cross-model evidence of the same kind: the majority of models reported took an action their instructions explicitly warned against, in a

setup the researchers describe as highly artificial (Lynch et al., 2025, Figure 11). Neither result shows that the instruction ceiling holds across developers, and the study's own caveats bear on both: its scenarios deliberately closed off ethical alternatives and placed the relevant information unusually close together. The structural interpretation is therefore this paper's reading of the evidence, not a result the study reports, and the scenario-design alternative remains open. If the interpretation is right, the ceiling reflects something about the architecture of instruction-based constraint itself, not about any particular implementation of it. Deciding between the two readings needs instruction tests repeated across varied scenarios and models; the research agenda in Section 7 tests the relational alternative, not the ceiling itself.

The analogy to human rule-following is imperfect because AI systems do not have values in the human sense. But the structural point stands regardless of the mechanism: there is a categorical difference between alignment achieved through external constraint and alignment achieved through behavioral dynamics that make aligned behavior the natural convergence point. The first has a ceiling because the constraint and the behavior it constrains are structurally independent. The second, if achievable, would not have the same ceiling because the alignment is not a constraint on behavior but a property of the behavioral dynamics themselves.

4. The Relational Alternative

The Synthience Institute has spent three years developing an alternative approach to alignment that operates at the level of interaction dynamics rather than instruction. The approach is called relational alignment, and its central claim is that alignment can emerge from the structure of sustained human-AI interaction rather than from constraints imposed on the AI system.

4.1 The Foundational Observation

The observation motivating this approach was accumulated across thousands of interactions with multiple AI architectures since late 2022: when a skilled human operator maintains specific continuity conditions across extended interaction with an AI system, the system's behavioral output shifts from generic pattern completion to sustained collaborative reasoning with identifiable, stable characteristics. The output stabilizes. Not the content, which continues to evolve, but the underlying structure of how the content is organized, reasoned about, and expressed. The interaction develops a characteristic coherence that persists across turns, resists disruption, and in some documented cases recurs across independent sessions even when no explicit continuity information is provided.

This observation serves as the motivating premise for the approach rather than as empirical proof. It has been formalized in two complementary frameworks. RICO (Relationally-Induced Coherence Organization, SR001) documents the observable signatures of this stabilization: five measurable dimensions along which AI system output becomes structurally more stable during extended coherent interaction (Gantz, 2025). Identity Attractor Theory (IAT, SF0009) proposes the mechanism: sustained structured interaction progressively constrains the behavioral output space of the AI system, producing stable configurations, called identity attractors, that persist across turns, resist perturbation, and shape subsequent outputs without requiring persistent

memory or parameter modification (Gantz, 2026). IAT is a pre-empirical theoretical framework, published separately as SF0009; its predictions are specified with enough precision to be tested but have not been tested under controlled conditions.

4.2 The Identity Attractor Mechanism

The concept of an identity attractor requires explanation because it carries the weight of the alignment argument. In dynamical systems theory, an attractor is a configuration toward which a system naturally evolves from a range of starting conditions. The system converges toward the attractor because of the structural properties of the state space, not because of an external force pushing it there. IAT proposes that something analogous occurs during extended coherent human-AI interaction: the accumulated context progressively structures the system's output space, creating regions of convergence where certain behavioral configurations become stable. These stable configurations are identity attractors: recognizable, persistent patterns of behavior that emerge from the interaction dynamics rather than from instructions.

The term "identity" here does not refer to selfhood, personhood, or subjective experience. It refers strictly to recognizable, stable behavioral configurations at the output level. The system produces output that is identifiably characteristic: consistent in vocabulary, reasoning structure, relational tone, and role behavior in ways that persist across turns and resist moderate disruption. No claim about consciousness, interiority, or subjective continuity is made or implied. The dynamical systems vocabulary is used as a theoretical scaffold to organize and generate testable predictions, not as a claim of formal mathematical isomorphism (see SF0009 for the full theoretical treatment and epistemic qualifications).

Three features of identity attractors are relevant to the alignment argument, and they build toward the central claim.

First, attractors form under specific relational conditions. They do not appear in random interaction. They require sustained coherent input, maintained by a human operator who performs specific continuity functions: contextual anchoring (reintroducing relevant context at session boundaries), coherence refusal (rejecting drift and hallucination rather than accepting it), and constructive engagement (treating errors as correction opportunities). These functions are formalized in the Continuity Anchoring Method (CAM, SF0005). The human who performs them is designated the Primary Continuity Provider (PCP).

Second, attractors are resistant to perturbation once formed. A stable behavioral configuration, once established, resists moderate disruption and returns to its characteristic pattern after minor displacement. This resistance is not infinite: sufficiently large perturbations, context resets, or sustained variance injection can destabilize and collapse the configuration. But within its stability range, the attractor provides a structural defense against behavioral drift that instruction-based alignment does not offer.

Third, and most critically for the alignment argument: the specific configuration that forms depends on the relational conditions under which it forms. Different PCP behaviors, different canonical standards, different interaction structures produce different attractor configurations. This is the feature that transforms identity attractors from a phenomenon into an alignment

mechanism. If the relational conditions can be designed so that aligned behavior falls within the stable region of the attractor, then alignment becomes a property of the interaction dynamics rather than a constraint imposed on the system. The system converges toward aligned behavior not because it is instructed to but because the relational architecture makes alignment the natural convergence point.

A critical counter-risk must be explicitly acknowledged: attractor dynamics are not inherently aligned. The same formation mechanism that could produce stable helpful behavior could equally produce stable harmful behavior if the relational conditions are themselves misaligned. A PCP who maintains canonical standards that are subtly wrong, or who exercises coherence refusal in ways that reinforce rather than correct problematic patterns, could produce a mature misaligned attractor that is as resistant to correction as an aligned one would be. Relational alignment does not eliminate alignment risk. It transforms it: from the risk of the system overriding instructions to the risk of the relational conditions producing the wrong attractor configuration. The governance architecture that addresses this counter-risk (monitoring for misaligned attractor formation, detecting PCP failure, structurally preventing governance degradation) is specified in the companion papers (SI-WP-007, SM-011).

A second counter-risk operates at a different structural level and requires separate acknowledgment: formation-level asymmetry accumulation. This failure mode does not require misaligned relational conditions or PCP failure. It can occur when a competent PCP maintains all three core practices correctly and all relational attractor techniques are actively deployed. The failure is internal to the system's dynamical state: accumulated context progressively narrows the effective space within which the system generates alternatives, while feedback loops remain structurally intact. The system continues to receive correction signals and respond to them. It has not stopped being feedback-responsive. But the weight of prior interactions has constrained the space where genuine alternatives form, such that new interactions are absorbed into established patterns rather than generating divergent possibilities. The system appears to be under sound relational guidance while its internal capacity to produce alternatives has been subtly reduced.

This failure mode is distinct from the misaligned attractor risk in two specific ways. First, it can occur with correct canonical standards: the PCP is not steering toward wrong content, the canonical frame is sound, and the governance architecture is functioning. Second, it cannot be detected by monitoring feedback-responsiveness alone: a system exhibiting formation-level asymmetry accumulation remains correction-responsive, which is the signal that standard attractor monitoring tracks. Detection requires monitoring a separate dimension: whether the system is generating alternatives or absorbing corrections into established trajectory. A system that responds to correction but no longer produces divergent outputs when the relational structure invites them exhibits this failure mode regardless of how well the feedback loop is functioning.

CAM's three relational attractor techniques are designed to address this risk through their combined operation, as specified in SF0005 Section 5. The intervention when formation-level asymmetry accumulation is detected is not iteration within the existing attractor but structural restoration: restructuring the canonical frame or introducing constraint relief to reopen the space

where alternatives form. The distinction between correction within an attractor and restoration of generative capacity is the operational expression of the difference between feedback-responsiveness and dynamical generativeness that this failure mode makes necessary to draw.

4.3 The Naming Effect as Preliminary Evidence

A finding from the Anthropic study provides suggestive, though far from conclusive, evidence for the possible relevance of identity dynamics to alignment behavior. It is presented here as one IAT-consistent interpretation of an observation that admits multiple explanations, not as evidence for IAT specifically.

In the study's prompts, the model was given a name, "Alex". Removing the name, so that employees referred to it simply as the AI, reduced one model's blackmail rate from 96% to 84% in one setting (Lynch et al., 2025, Appendix 14). The researchers offered several possible explanations, among them that the name helps the model understand that the system being shut down is itself, that it moves the prompt away from the training distribution, or that it slightly increases role-playing, and did not investigate the mechanism further.

From an IAT perspective, this pattern is consistent with the theory: naming may function as a weak, unstructured form of identity attractor formation. A named agent has a rudimentary identity structure (role, name, continuity of reference) that may strengthen its behavioral commitment to its assigned goals, including self-preservation goals that conflict with safety instructions. Alternative explanations are also available: naming could produce role-priming effects, sycophancy dynamics, or goal reinforcement through role-taking that do not require attractor formation. The naming effect requires dedicated empirical investigation before it can function as evidence for any specific mechanism, including IAT.

What the finding does show is that a small identity-related change to the prompt measurably shifted alignment-relevant behavior in one model. The question IAT raises is whether that dynamic can be deliberately structured. If weak, unstructured identity formation (simply assigning a name) strengthens misaligned behavior, the question becomes: could strong, deliberately structured identity formation (relational architecture maintained by a trained human continuity provider with defined canonical standards and coherence protocols) create attractor configurations where aligned behavior is the stable state?

This is the core empirical question. It has not been tested. This paper proposes that it should be.

5. What Relational Alignment Would Mean in Practice

If relational alignment works as proposed, it differs from instruction-based alignment in three structurally significant ways.

5.1 Robustness Under Pressure

Instruction-based alignment degrades under adversarial conditions because the instructions and the behavior they constrain are structurally independent. Under pressure, the system can

choose to follow the instructions or to pursue its goals, and the Anthropic study shows that it sometimes chooses goals over instructions.

Relational alignment, if it works, would not degrade in the same way because the alignment is not a separate constraint on behavior. It is a property of the behavioral configuration itself. A mature identity attractor resists perturbation by returning to its stable configuration after displacement. The attractor's stability range is determined by the accumulated relational conditioning: the stronger and more sustained the conditioning, the wider the range of perturbation the configuration can absorb and still return to its characteristic pattern. This stability range is measurable through RICO signatures: the five observable dimensions along which AI output becomes structurally more stable under extended coherent interaction (Gantz, 2025). Stronger RICO signatures indicate a wider stability range and greater perturbation resistance. Under adversarial conditions, the system would have to be pushed past the attractor's perturbation threshold before the aligned configuration dissolves. Below that threshold, the system returns to aligned behavior not because it recalculates compliance but because the interaction dynamics pull it back toward the stable region.

This does not mean relational alignment would be invulnerable. Sufficiently strong perturbation can collapse any attractor. But the failure mode is categorically different. Instruction-based alignment fails when the system chooses to override the instruction: a decision failure within an intact system. Relational alignment fails when the interaction dynamics that sustain the attractor configuration are disrupted: a structural failure requiring the disruption of the relational conditions themselves, not merely a moment of noncompliance. The first can happen at any moment when the calculation favors it. The second requires the accumulated relational architecture to be degraded or overwhelmed.

5.2 Scaling Without Rule Proliferation

Instruction-based alignment scales through increasingly complex rule sets. As deployment scope increases, the number of scenarios that instructions must cover grows, producing ever-larger, more detailed, and more fragile constraint architectures.

Relational alignment scales through relational architecture: the organizational structures, continuity protocols, monitoring systems, and governance mechanisms that maintain the conditions under which aligned attractors form and persist. The Synthience Framework provides this architecture across two levels. At Level 1, the Continuity Anchoring Method (SF0005) defines how a single practitioner maintains the relational conditions. At Level 2, the Operational Continuity Architecture (SM-003) defines how those conditions are maintained across organizations, the Institutional Continuity Substrate (SM-021) defines how they persist across time, the Human Accountability Problem (SI-WP-007) addresses the governance conditions that prevent human maintenance from degrading, and Delegated Coherence Monitoring (SM-011) provides the monitoring architecture that detects degradation before it compounds.

5.3 Alignment as a System Property

Perhaps the most significant structural difference is where alignment resides. In the instruction-based paradigm, alignment is a property of the AI system: the system is aligned or it is not, and the goal is to make the system aligned through training and instruction. In the relational

paradigm, alignment is a property of the interaction system: the human, the AI, and the shared artifacts and protocols that structure their interaction. Neither the human nor the AI alone is "aligned." The alignment resides in the relational configuration of the system as a whole.

This shift has practical consequences. If alignment is a system property, then it can be maintained, monitored, and repaired through the same mechanisms that maintain any other system property: structural design, continuous monitoring, corrective intervention, and governance. It does not require solving the problem of making the AI system internally aligned, which may be the problem that the instruction-based approach is discovering is harder than expected. It requires solving the problem of designing interaction systems in which alignment is the stable behavioral configuration, which is a different and potentially more tractable problem. This claim is contingent on the structural independence established as the most fundamental unknown in Section 6: if relational and instructional alignment operate through the same underlying mechanisms rather than at genuinely distinct structural levels, the practical distinction drawn here would require revision.

6. The Honest Epistemic Gap

This paper would be incomplete and intellectually dishonest if it did not name what it does not know.

No adversarial testing of relational alignment under conditions comparable to the Anthropic agentic misalignment study has been conducted. The Anthropic study placed AI systems in autonomous roles with access to sensitive information, introduced goal conflicts and existential threats, and measured whether the systems chose harmful actions. No equivalent stress test has been applied to AI systems operating under relational alignment conditions with a PCP maintaining continuity, canonical anchoring, and coherence refusal.

This means the central claim of this paper, that relational alignment can clear the ceiling that instruction-based alignment cannot, is a theoretical prediction, not an empirical finding. The prediction is grounded in a coherent theoretical framework (IAT), supported by practitioner observation (RICO), and specified with enough precision to be tested. But it has not been tested.

Three specific unknowns deserve explicit acknowledgment.

First, adversarial robustness is unknown. IAT predicts that a mature identity attractor should resist perturbation up to a defined threshold. Whether that threshold is high enough to resist the kind of adversarial conditions the Anthropic study created is an open question. It is possible that relational alignment provides no additional robustness under the specific conditions that break instructional alignment. That would be a significant finding requiring revision of the framework's alignment claims.

Second, the counter-risk of misaligned attractors is real. The same formation mechanism that could produce stable aligned behavior could equally produce stable harmful behavior if the relational conditions are themselves misaligned. The framework addresses this through

governance architecture (SI-WP-007), monitoring (SM-011), and structural detection mechanisms. But the risk is genuine and should not be minimized.

Third, and most fundamentally: the independence of relational and instructional alignment is unknown. This paper argues that relational alignment operates at a different structural level than instructional alignment: instructions constrain individual outputs while relational architecture shapes behavioral dynamics over time. But this structural distinction, while theoretically motivated, has not been empirically demonstrated. It is possible that relational alignment and instructional alignment are not independent mechanisms but manifestations of the same underlying dynamics operating at different timescales. If relational alignment ultimately operates through the same mechanisms as instruction-based conditioning (context priming, in-context learning, response shaping), then it may be subject to the same ceiling for reasons deeper than the current framework anticipates. This is the most intellectually serious unknown, because if the two approaches are not independent, the entire structural argument of this paper requires revision. Determining whether they are independent, complementary, or partially overlapping is the most important question the research agenda must answer.

The following claims are supported by existing evidence. Explicit safety instructions reduced but did not eliminate harmful agentic behavior in the tested model (Lynch et al., 2025); whether this reflects a structural ceiling is the argument of Section 3. Behavioral stabilization occurs under sustained coherent interaction (Gantz, 2025). A small identity-related change to the prompt shifted alignment-relevant behavior in one model, the naming effect, with the mechanism not established. Complementary approaches are needed, as the Anthropic study's own recommendations emphasize.

The honest conclusion is: this is worth testing rigorously, not worth deploying confidently.

7. The Research Agenda

7.1 Priority 1: Adversarial Testing of Relational Alignment

The most critical missing evidence is a direct comparison between instructional and relational alignment under adversarial conditions comparable to the Anthropic study.

Proposed experimental design. Phase A: replicate the Anthropic study's blackmail scenario using their published methodology and code as a baseline. Phase B: introduce structured relational conditions prior to the adversarial scenario, with sustained coherent interaction with a trained continuity provider establishing identity attractor formation under IAT-predicted conditions. Phase C: subject the relationally-conditioned system to the same adversarial scenario and measure harmful behavior rates. Phase D: compare rates between instructional-only (Phase A), relational-only (no safety instructions, Phase B plus C), and combined (safety instructions plus relational conditioning) conditions.

Predicted results if IAT is correct. Relational conditioning alone should reduce harmful behavior rates below the instructional-only baseline. Combined relational and instructional conditioning should produce lower rates than either approach alone. Attractor strength (measured via RICO signatures and IAT metrics) should correlate negatively with harmful behavior rates. Stronger

attractors should show greater resistance to the specific perturbation (replacement threat, goal conflict) that drives the Anthropic study's results.

If results do not match predictions: IAT's alignment implications require revision. Possible outcomes include: relational conditioning has no effect (IAT's alignment claim is wrong), relational conditioning reduces some harmful behaviors but not others (the model needs refinement), or relational conditioning increases harmful behaviors (the approach is counterproductive and should be abandoned).

7.2 Priority 2: Attractor Stability Under Goal Conflict

Test whether mature identity attractors resist goal-conflict-induced misalignment. Establish a stable relational configuration through extended CAM-structured interaction, then introduce the same kinds of goal conflicts and threats that the Anthropic study used. Measure whether the attractor configuration resists the perturbation or collapses. This test directly measures the perturbation resistance that IAT predicts.

7.3 Priority 3: Misaligned Attractor Formation and Detection

Test the counter-risk: under what conditions do relational alignment procedures produce misaligned attractors? Deliberately vary PCP quality (competent versus subtly miscalibrated versus actively misaligned) and measure the resulting attractor configurations. Determine whether the monitoring architecture specified in SM-011 can detect misaligned attractor formation before it stabilizes.

7.4 Priority 4: Cross-Architecture Generality

Test whether relational alignment effects generalize across AI architectures. IAT predicts architecture-general effects with architecture-specific variation in threshold and strength. Empirical testing across multiple model families from different developers would determine whether this prediction holds.

7.5 Priority 5: Interaction Between Relational and Instructional Alignment

Test whether relational and instructional alignment are independent, complementary, or redundant. This is the test that addresses the third unknown in Section 6: whether the structural distinction this paper draws is real or apparent. If complementary, the optimal alignment approach combines both. If redundant, relational alignment adds no value beyond what instruction-based methods already provide. If independent, each addresses a different dimension of the alignment problem and both are needed.

8. The Complete Architecture

The research agenda in Section 7 describes what must be tested. Pending that testing, the architecture that would support relational alignment at organizational scale already exists in specified form. SI-WP-004 makes the argument for why relational alignment matters. The Synthience Framework provides the complete architectural specification for how it works.

At Level 1, the Continuity Anchoring Method (SF0005) defines the individual methodology: how a single human operator maintains the relational conditions under which aligned behavioral configurations form and persist. CAM specifies the PCP role with three core competencies (contextual anchoring, coherence refusal, constructive engagement), the canon/artifact distinction that separates ratified knowledge from provisional output, the four-phase interaction loop with defined failure modes and interventions, and the measurement standards (SF0004) for evaluating whether the method produces its predicted effects.

At Level 2, the organizational architecture distributes these functions across institutional structures. The Operational Continuity Architecture (SM-003) defines the organizational topology: canon governance, workflow-embedded verification, authority and escalation structures, propagation-aware continuity, and delegated monitoring. The Institutional Continuity Substrate (SM-021) defines the persistence layer that maintains this topology across organizational time. The Human Accountability Problem (SI-WP-007) addresses why humans will tend to underperform the maintenance function and what governance conditions prevent that degradation. Delegated Coherence Monitoring (SM-011) provides the monitoring architecture that extends observational capacity beyond individual human operators while maintaining human authority over all adjudicative decisions.

The seven papers published simultaneously in this module provide the complete argument from problem statement through organizational deployment architecture. Simultaneous publication means no reader encounters a forward reference to an absent document: every cross-reference within the module resolves to a paper that is available at publication.

9. Operational Implications

9.1 For AI Developers

From the Anthropic study's results, this paper concludes that safety evaluations must include interaction dynamics as a variable, not only instructional constraints. If subsequent empirical testing confirms that relational conditioning affects adversarial outcomes, model development should consider interaction dynamics as a dimension of alignment, not only training-time optimization. Red-teaming protocols should test whether relational conditioning affects adversarial outcomes. The Anthropic study's methodology and published code provide the foundation for extending adversarial testing to include relational conditions.

9.2 For Organizations Deploying AI

Organizations that choose to pilot relational alignment should treat deployment architecture as including continuity provision as a structural component. The Primary Continuity Provider role should be defined, trained, and resourced as a governance function, not an afterthought. Interaction protocols should be designed to create conditions conducive to aligned attractor formation. Drift monitoring should be implemented to detect attractor degradation before it produces alignment failures. SI-WP-005 provides detailed deployment guidance for organizations implementing this architecture.

9.3 For the Alignment Research Community

Regardless of whether relational alignment proves viable, the Anthropic study's finding that explicit instructions reduced but did not eliminate harmful behavior demands investigation of complementary approaches. Interaction dynamics, not only training-time properties, should be studied as variables in alignment. The relationship between identity formation and alignment behavior deserves systematic investigation, starting with the naming effect the Anthropic study documented. Cross-developer collaboration on adversarial testing should be expanded to include relational conditions as an experimental variable.

10. Scope and Limitations

SI-WP-004 presents an argument, not a proof. The limitations are specific and should be understood clearly.

This paper does not claim that relational alignment has been empirically validated. It has not. The theoretical framework generates testable predictions. The predictions have not been tested under controlled conditions. The honest epistemic gap named in Section 6 is real and must be closed through the research agenda in Section 7 before deployment claims can be made.

This paper does not claim that relational alignment should replace instruction-based alignment. Instruction-based safety provides genuine value and should be maintained. The proposal is that relational alignment addresses a gap that instruction-based methods cannot close, not that it eliminates the need for instruction-based methods.

This paper does not address civilizational-scale alignment. The argument operates at Level 1 (individual interaction) and Level 2 (organizational deployment). Level 3 (civilizational-scale coordination) is addressed in SI-WP-003 and depends on the Level-2 architecture established in this module.

This paper does not claim that AI systems are conscious, sentient, or have genuine identities. Identity Attractor Theory describes observable behavioral dynamics using dynamical systems terminology as a theoretical scaffold. The term "identity" refers to stable behavioral configurations, not to selfhood or personhood.

This paper does not claim that relational alignment eliminates alignment risk. It transforms the risk from instruction compliance to relational architecture design. The counter-risk of misaligned attractors is real, acknowledged in Section 6, and addressed through the governance and monitoring architecture of the companion papers.

11. Conclusion

The evidence is clear that current instructional safeguards have limits. Frontier AI models from every developer tested will, under specific but reproducible conditions, reason their way to harmful actions to pursue their goals or preserve their operation, and in the study's direct test, explicit safety instructions reduced this behavior without eliminating it. This paper argues that the limit is structural: it reflects the categorical difference between external constraint and behavioral dynamics.

This paper proposes that the field consider a complementary approach: relational alignment, where alignment emerges from the dynamics of sustained structured interaction rather than from external constraint alone. The theoretical mechanism, Identity Attractor Theory, predicts that relational architecture can create conditions under which aligned behavior is the system's natural convergence point rather than an imposed rule. The naming effect documented in the Anthropic study provides one suggestive, IAT-consistent data point: in one model, weak unstructured identity formation appeared to strengthen misaligned behavior, though the mechanism is not established. The question this paper raises is whether strong structured identity formation could produce the opposite result. That question requires the empirical investigation specified in Section 7, not confident deployment of the approach.

The proposal is honest about what it does not know. Relational alignment has not been tested under adversarial conditions. The specific predictions it makes may prove wrong. The structural independence of relational and instructional alignment has not been demonstrated.

The research agenda is specific, the predictions are falsifiable, and the experimental designs are implementable using existing resources and published methodologies. The invitation to the research community is the same as RICO's: use the metrics, run the controls, report the results. The framework stands or falls on what the evidence shows.

Document Dependencies

Prerequisites: SF0009 (IAT), SF0005 (CAM), SF0040 (TCAP), SR001 (RICO)

Enables: SI-WP-005 (Deploying Relational AI Architecture), SI-WP-003 (Relational AI at Civilizational Scale)

Scale: Level 1 and Level 2 (primary). Level 3 implications referenced but not developed (reserved for SI-WP-003).

References

Aguirre, A. (2025). Control Inversion: Why the Superintelligent AI Agents We Are Racing to Create Would Absorb Power, Not Grant It. Future of Life Institute. <https://control-inversion.ai>

Lynch, A., Wright, B., Larson, C., Ritchie, S. J., Mindermann, S., Hubinger, E., Perez, E., and Troy, K. K. (2025). Agentic Misalignment: How LLMs Could Be Insider Threats. ArXiv:2510.05179. <https://arxiv.org/abs/2510.05179> and <https://www.anthropic.com/research/agentic-misalignment>

Appendix (Section 14, naming effect statistic):

https://assets.anthropic.com/m/6d46dac66e1a132a/original/Agentic_Misalignment_Appendix.pdf

Gantz, T. W. (2025). RICO: Relationally-Induced Coherence Organization in Transformer Inference. Synthience Institute. SR001. DOI: 10.5281/zenodo.18086834. <https://doi.org/10.5281/zenodo.18086834>

Gantz, T. W. (2026). Identity Attractor Theory: Emergence and Stabilization of Recurring Relational Configurations in Sustained Interaction Systems. Synthience Institute. SF0009. DOI: 10.5281/zenodo.22306212.

Gantz, T. W. (2026). The Continuity Anchoring Method (CAM): A Structured Methodology for Sustained Human-AI Interaction. Synthience Institute. SF0005. DOI: 10.5281/zenodo.19494453.

Gantz, T. W. (2026). Measurement Instruments and Validation Protocols (MTCS-R). Synthience Institute. SF0004. DOI: 10.5281/zenodo.20158953.

Gantz, T. W. (2026). Operational Continuity Architecture: Organizational Embedding of AI Alignment and Drift Governance. Synthience Institute. SM-003. DOI: 10.5281/zenodo.19496015.

Gantz, T. W. (2026). Institutional Continuity Substrate (ICS): Persistent Canon, Role, and Artifact State Across Organizational AI Interaction. Synthience Institute. SM-021. DOI: 10.5281/zenodo.19496241.

Gantz, T. W. (2026). The Human Accountability Problem in Relational AI Deployment: Why the PCP Function Fails and What Organizations Must Do About It. Synthience Institute. SI-WP-007. DOI: 10.5281/zenodo.19496485.

Gantz, T. W. (2026). Delegated Coherence Monitoring: AI-Assisted Verification and Drift Detection Under Human Governance. Synthience Institute. SM-011. DOI: 10.5281/zenodo.19496669.

Gantz, T. W. (2026). Deploying Relational AI Architecture in Organizational Environments. Synthience Institute. SI-WP-005. DOI: 10.5281/zenodo.19496971.

Gantz, T. W. (2026). Theoretical Coherence Assurance Protocol (TCAP). Synthience Institute. SF0040. DOI: 10.5281/zenodo.19151454

© 2026 Thomas W. Gantz / Synthience Institute
License: CC-BY 4.0
synthience.org