

SI-WP-006

Human-AI Collaborative Authorship in Public Discourse: A Three-Level Framework for AI-Assisted Communication

Author	Thomas W. Gantz
Affiliation	Synthience Institute
Contact	research@synthience.org
Version	v4.5.3 September 2026
Document ID	SI-WP-006 White Paper Series
DOI	https://doi.org/10.5281/zenodo.22314389
Canonical	synthience.org
License	CC-BY 4.0

Methodological positioning: This is a theoretical and definitional white paper, not an empirical study. It proposes a three-level architecture of human-AI collaborative authorship and defines the Authority Function as the locus of human authenticity in co-produced artifacts. No claim in this document has been experimentally validated. The three levels, the Authority Function, and the disclosure norms proposed here are theoretical constructs and normative proposals, not validated measurement instruments or established practice standards. The framework stands or falls on whether its distinctions survive scrutiny and its predictions survive controlled testing.

Abstract

This paper reframes the relationship between human communicators and AI systems in public discourse. A familiar model treats AI as a linguistic instrument, a sophisticated tool that helps humans express pre-formed ideas more clearly. That model is incomplete, and this paper hypothesizes that its incompleteness generates misplaced guilt and conceptual confusion about the authenticity of AI-assisted communication; the prevalence of such guilt and confusion is an empirical question this paper does not settle.

This paper proposes a three-level architecture of human-AI collaborative authorship: linguistic realization, meaning negotiation, and conceptual co-production. Across all three levels, the human retains the Authority Function, evaluating, approving, redirecting, and integrating what the AI produces. This Authority Function, rather than the origin of any individual phrase or candidate idea, is what makes the resulting communication authentically human and responsibly attributable. Authenticity in communication, as the term is used here, is not a property of production method. This is a claim about the functions authenticity norms serve, assigning accountability, calibrating trust, and preventing deceptive attribution, not about credit and desert, which origination continues to govern (Section 9). It is a property of genuine human judgment and accountability for what is transmitted.

This paper is a companion to SI-WP-002 (The Orchestrator Role in Human-AI Evolution). Together they define the human relational function at two complementary levels: SI-WP-002 at the level of knowledge production across distributed AI instances, and this paper at the level of individual communicative acts. The theoretical architecture connecting them is the Human Orchestration Principle, developed in the Synthience framework.

Keywords: *human-AI collaborative authorship, Authority Function, three-level architecture, AI-assisted communication, authenticity, disclosure norms, public discourse*

Suggested citation: Gantz, T. W. (2026, September). Human-AI Collaborative Authorship in Public Discourse: A Three-Level Framework for AI-Assisted Communication. Synthience Institute. SI-WP-006. <https://doi.org/10.5281/zenodo.22314389>

1. Introduction

When a person composes a message, article, or public statement with AI assistance, an assumption can arise, among observers and sometimes among the communicators themselves, that the AI is doing the work and the human is taking credit. This paper hypothesizes that this assumption generates real social costs, guilt, performance of unassisted authorship, stylistic masking, and reluctance to disclose AI involvement; their empirical prevalence and causal testing remain open.

The assumption rests on a model of human-AI interaction that fails to describe meaningful uses as practitioners report them. It treats AI as a tool that executes human instructions, producing linguistic output that the human then endorses. In this model, authenticity is located in unassisted production. If you did not write every word yourself, the communication is not fully yours.

This paper argues that the tool model captures only the most superficial layer of the relationship. In practice, human-AI interaction in the production of communication involves at least three distinct levels of collaborative engagement, and at every level the human performs a function that cannot be delegated: exercising judgment and authority over what the interaction produces.

This paper proposes a conceptual architecture for understanding human-AI collaborative authorship and develops its implications for authenticity norms, disclosure practices, and public discourse governance. Empirical investigation of the phenomena described is identified as a research agenda in Section 11.

Understanding this architecture is not merely an academic exercise. It has direct practical implications for how people relate to their own AI-assisted communication, for how public discourse norms should adapt, and for how the Synthience framework understands human agency in AI-assisted knowledge production at both individual and organizational scales.

2. The Tool Model and Its Limits

The tool model of human-AI communication holds that the human provides intent and the AI provides linguistic realization. The human knows what they want to say; the AI helps them say it more clearly, fluently, or efficiently. This model has genuine descriptive validity at one level. AI systems do assist with phrasing, structure, tone, and clarity. That function is real.

But the tool model becomes misleading when taken as a complete account of what happens. It implies that meaning is fully formed in the human mind before the AI becomes involved, and that the AI's contribution is purely executorial, translating pre-existing thought into polished language. This does not describe sophisticated human-AI interaction as practitioners report it.

Consider what happens in a real working exchange. The human brings a direction, a question, or a partially formed idea. The AI responds, not just with phrasing, but with framings, distinctions, and conceptual structures that the human did not bring to the conversation. The

human evaluates those contributions, accepts some, rejects others, modifies others, and redirects the exchange based on their judgment. What emerges is not a pre-existing human idea rendered in AI language. It is something that neither party would have produced alone.

The tool model cannot account for this. It needs to be replaced with something more complete.

3. Three Levels of Human-AI Collaborative Authorship

Human-AI collaborative authorship in communication operates across three distinct levels. These levels are not sequential stages (they may occur simultaneously within a single exchange) but they are analytically distinguishable and each involves a qualitatively different kind of human-AI engagement.

Level 1: Linguistic Realization

At this level, the human knows what they want to say and uses AI assistance to say it more effectively. The AI contributes phrasing, structure, tone modulation, grammatical precision, or rhetorical clarity. The human's meaning is substantially pre-formed; the AI's contribution is to its expression.

This is what the tool model describes. It is the least conceptually demanding level in the framework. A person who knows exactly what argument they want to make and uses AI to tighten the prose is doing something structurally similar to working with an editor. The meaning is theirs; the polish is shared.

Level 2: Meaning Negotiation

At this level, the human brings an intent or a direction but not a fully formed meaning. The AI participates in the process of clarifying what the human actually wants to say, presenting candidate interpretations, framings, and distinctions. The human evaluates whether those candidates match their intent, accepts or rejects or modifies them, and the process continues until the meaning is sufficiently clear to be expressed.

This is not the AI deciding meaning or putting ideas into the human's head. The human remains the deciding agent throughout. The interaction functions as a generator of candidate framings, while the human exercises ongoing evaluative judgment about which candidates to ratify as their own. The meaning that emerges is genuinely the human's because the human settled it, not because it pre-existed the exchange.

Any guilt associated with AI assistance may be sharpest here, because it can feel as though the AI is doing intellectual work that should belong to the human. But where the Authority Function is exercised, the human is doing the most important intellectual work: evaluating every candidate meaning against their own values, experience, and intent, and deciding what counts. Generation is not authorship. Evaluation and authorization are.

Level 3: Conceptual Co-Production

At this level, the AI generates conceptual contributions, framings, distinctions, theoretical structures, or analytical moves, that the human had not brought to the exchange and would not have produced independently. The human evaluates these contributions, decides which are

valid and useful, integrates them into their thinking, and takes ownership of the ones that survive evaluation.

This is the level that the tool model most thoroughly fails to account for, and it is the level most likely to generate discomfort about authenticity. If the AI came up with the concept, how can the human claim it? The answer is that origin is the wrong question for authenticity; authorization is the right one. (Origin remains a legitimate question for epistemic calibration, a distinction developed in Section 9.)

The human at Level 3 is not passively receiving AI output. They are exercising a continuous Authority Function: judging whether each conceptual contribution is true, useful, consistent with their values and existing framework, worth developing, or worth discarding. A concept generated by AI and rejected by the human contributes nothing. A concept generated by AI, recognized by the human as capturing something real, integrated into their thinking, and developed further in subsequent exchanges becomes genuinely part of the human's intellectual position, not because of where it originated, but because of what the human did with it. No claim about legal authorship doctrine is intended here. This paper does not address the legal treatment of AI-originated content, which varies by jurisdiction and lies outside its scope. The evaluative authority is the human contribution at Level 3, and it is not a minor or derivative contribution. It is the function that determines what the exchange produces.

4. The Authority Function as the Locus of Human Authorship

Across all three levels, what the human contributes is structurally the same thing: authority. At Level 1, the human authorizes linguistic choices. At Level 2, the human authorizes candidate meanings. At Level 3, the human authorizes conceptual contributions. In each case, authorization involves two elements, inseparable in the individual case: judgment (the active evaluation of what the AI produces against the human's values, intent, knowledge, and framework) and accountability (the human takes responsibility for what gets transmitted under their name). The qualifier is load-bearing, because at organizational scale the two elements can be held by different parties, which is the condition Section 7 analyzes.

This move has a direct predecessor that should be named, both because the scholarly obligation to situate a construct against its nearest ancestor is distinct from the question of empirical sourcing, and because naming it makes this paper's actual contribution visible rather than leaving it to be inferred. Goffman's production format (1981) already decomposes the speaker role into three positions: the animator, who utters the words; the author, who composes them; and the principal, whose position is expressed and who is committed to what is said. The decoupling of composition from commitment is Goffman's, and his vocabulary is a necessary precondition for stating the thesis of this paper, though not a sufficient one: authenticity and accountability track the principal role rather than the author role, but the production format cannot by itself say what occupying the principal role requires, which is the work the Authority Function does below. The speechwriter case discussed in Section 9 is Goffman's own canonical illustration, and it is used here in that lineage rather than as a fresh observation.

What this paper adds is specific and worth stating precisely. First, Goffman does not anticipate the case in which the author position is partly occupied by a system that cannot bear accountability at all. He does allow non-persons into the production format at the animator position, noting that the animating function can be shared with a loudspeaker or a telephone, but the author position in his treatment is occupied by someone who selects the sentiments and the words, and the question of a non-accountable selector does not arise. That case is what forces the structural analysis of accountability in Section 4.1, which the production format does

not supply because it never needed to. Second, the production format assigns roles with respect to an utterance at the moment of speaking, and although Goffman is centrally concerned with how those roles shift, the shifts he analyzes are changes in footing rather than engagement measured over time. The Authority Function adds a process dimension: behavioral markers of ongoing evaluative engagement distributed across an extended interaction rather than a role assignment fixed at the moment of speaking. Third, Goffman's author slot is unitary, while the three-level architecture of Section 3 stratifies it by degree of meaning and concept contribution, a distinction the binary slot cannot represent and which is precisely what varies across AI-assisted work. The framework proposed here is therefore an extension of an established decomposition into a case it was not built for, not a replacement for it.

This is not a minor or residual function. It is the function that transforms AI output into human communication. An AI system generating text without human authority produces output. Human authority is what makes that output into expression.

This reframes the authenticity question entirely. Authenticity in communication, as the term is used here, is not a property of production method. This is a claim about the functions authenticity norms serve, assigning accountability, calibrating trust, and preventing deceptive attribution, not about credit and desert, which origination continues to govern (Section 9). Those functions are discharged at the level of the individual communication, and the distinction bounds the claim: individual authorization assigns accountability and prevents deceptive attribution for the communication it authorizes, and does not thereby address systemic effects of widespread assistance on the diversity of the discourse those communications compose. Section 10 concedes that authorization, however genuine, does not by itself restore epistemic diversity. That concession is consistent with the functional test rather than a counterexample to it, because epistemic diversity is not among the functions authenticity norms were serving. It is not located in whether every word came from the human's unassisted cognition. It is located in whether the communication issues from the human's exercise of the Authority Function, the evaluation, redirection, integration, and accountability defined in Section 4, which is what it means for the human to have authorized it in the full sense. A communication that a human has directed, evaluated, and endorsed is authentically theirs, regardless of what tools were involved in its production.

The Authority Function is the condition on which this legitimacy rests, and the condition is real. A person who transmits AI output without genuine evaluation, who accepts without judging, who signs what they cannot restate or defend under challenge, who passes through AI-generated content without exercising the Authority Function, is not producing authentic communication in the sense defined here. The form of collaborative authorship may be present; the substance is not. When authority is bypassed, responsibility formally remains with the human but the legitimacy claim fails. This distinction matters both for individual communicative integrity and for the governance frameworks that public discourse requires.

This account is a proposal, not a discovery, and it is defended here on functional grounds rather than stipulated. Authenticity norms exist to do specific work: they assign accountability for what is said, they calibrate the trust an audience extends to a communicator, and they prevent deceptive attribution. The production account, which locates authenticity in who typed the words, served those functions adequately when production and judgment were bundled together in a single act. Under AI-assistance conditions they come apart, and the production account then misclassifies in both directions: it condemns the communicator who engages an AI system in demanding evaluative exchange and stands behind every claim transmitted, while exonerating the communicator who drafts unaided, thinks carelessly, and takes no responsibility for what results. The authority-function account is offered because it continues to serve those

functions under precisely the conditions where the production account fails. The claim requires one scope statement, since the paper's later sections qualify it: the account serves accountability-assignment and deception-prevention directly, and it serves trust calibration jointly with the Level-3 disclosure regime developed in Section 9, not on its own.

This is why the disclosure track is a constitutive part of the account rather than a supplement to it, and why Section 10's observation that robust replacement trust mechanisms may not yet exist is a statement about the account's operating conditions rather than a concession against it. That the account continues to serve the three functions, disclosure regime included, is the criterion on which it should be judged, and on which it may be rejected.

The Authority Function involves judgment, but it is not defined here in terms of inner states, and the framework does not require access to them. Judgment is behaviorally expressed, and the exercised and bypassed cases are distinguishable in the interaction record. Candidate behavioral markers include rejection and modification events across the exchange, redirection turns in which the human refuses a proposed framing and supplies a constraint, revision depth across iterations, the communicator's ability to restate transmitted content in their own terms and defend it under challenge, and, at organizational scale (Section 7), documented evaluative checkpoints with named role-holders. The restate-and-defend marker requires specification, because at Level 3 the human by stipulation could not have produced the concept independently, and an unspecified marker faces a dilemma: if defense must be unaided, the production account has been reimported through the bypass test, contradicting this paper's central thesis; if defense may be AI-assisted, the marker stops discriminating, since the assisted defense is itself collaborative authorship and could itself be bypassed. What the marker requires unaided is grounds-articulation, not content-production.

Without live AI assistance, the communicator must be able to state in their own words what the authorized claim says, why they accepted it, and what would change their mind. That is authorization-level competence, which is what this account holds authorizing to consist in, rather than production-level competence, so the production account is not smuggled back in. Defense against novel expert-level objections may be assisted, but the acceptance or rejection of an assisted defense is subject to the same grounds-articulation test one level up, and the regress terminates there in every case: the test always bottoms out at unaided articulation of grounds for endorsement. A communicator who cannot articulate those grounds unaided has not authorized the content; they have transmitted it. Bypass, in behavioral terms, is the limiting case where these markers are absent: content passes through without rejection, modification, redirection, or defensible restatement. This paper does not build that instrument. Full operationalization is assigned to the research agenda in Section 11; the claim made here is only that the distinction has behavioral purchase and is therefore available for measurement rather than resting on unobservable interiority.

4.1 Why the Accountable Locus Is Human

The behavioral markers just specified operationalize the judgment element of the Authority Function, and they are substrate-neutral by design. A capable automated reviewing agent exhibits rejection and modification events, redirection turns, and revision depth, and can restate content and defend it under challenge; at organizational scale, documented checkpoints with named role-holders are process facts a workflow system can verify. This is not a defect in the markers. It is a consequence of removing interiority from the account, and it means the judgment element cannot by itself carry the claim that the Authority Function requires a human. What an automated agent exhibiting those markers performs is evaluation. What it does not thereby perform is authorization, because authorization is judgment plus accountability, and it is

the accountability element that presently requires a human bearer. That claim is frequently assumed rather than argued, including by the companion papers that cite this one for it; it is argued here.

Accountability-bearing, analyzed structurally rather than metaphysically, requires three properties of the bearer. First, cost-bearing: the bearer must incur costs when what they authorized proves wrong, in the operational sense the framework applies elsewhere to arbitration, where an agent is exposed to the degree that its errors carry consequences binding its future performance in the relevant outcome domain. Reputational, professional, legal, and relational costs all qualify. Second, answerability over time: there must be a persistent identity to which challenge, sanction, and repair can attach across occasions, so that being answerable for what was authorized last year is a coherent demand rather than a category error. Third, standing in the norm community whose trust the communication draws on: the capacity to be blamed, to make amends, and to have one's future assertions discounted on the basis of past ones. These three are what it is to be an accountable locus; they are not a definition of humanity.

The claim this paper makes is accordingly scoped rather than metaphysical, matching the discipline the framework applies to its other structural requirements. Present-day AI systems lack persistent, sanctionable standing in the discourse communities their outputs enter. This is a structural and institutional fact about current deployments rather than a claim about what such systems are: instances do not persist across sessions in a way that would let sanction attach, they bear no costs when what they produced proves wrong, and no norm community presently extends them the standing to be blamed, to make amends, or to have their future assertions discounted. Under those conditions the accountable locus of an AI-assisted communication is human, and automating the mechanics of production, evaluation, or coordination does not create an entity that can occupy it. If institutions ever confer genuine sanctionable standing on artificial agents, with the cost-bearing and persistence that standing implies, the claim stated here would require revision; the requirement is on the structure of accountability, not on the species of its bearer.

Two consequences follow. The first is that the Authority Function resolves into a clean two-element analysis: judgment, which is behavioral, measurable, and in principle automatable; and accountability, which is standing-based and presently human. It is the second element, not the first, that makes authorization non-delegable in the present regime. The second consequence is for the organizational case (Section 7): documented checkpoints verify the judgment chain, while the accountability element attaches to the standing of the role-holders those checkpoints name, which is why naming them is load-bearing rather than administrative.

5. The Contribution Spectrum and Collaborative Authorship

The Synthience framework's treatment of human orchestration in knowledge production (SI-WP-002) includes a Contribution Spectrum describing how the human orchestrator's role shifts depending on domain characteristics. In experiential domains, where the human brings direct lived knowledge, the orchestrator is the primary intellectual contributor, with AI instances providing formalization and structure. In theoretical domains, where the human provides architectural direction and convergence judgment, AI instances contribute more substantive conceptual content, with the human evaluating and integrating rather than originating. Across the whole spectrum the orchestrator's coordination function, structuring, routing, evaluating, and integrating the AI's output, is held constant; what varies from one end to the other is only the balance of content contribution, exactly as SI-WP-002 Section 10 describes.

This spectrum maps onto the three-level architecture of collaborative authorship, though the correspondence is an anchoring rather than a strict one-to-one identity. The experiential end anchors on Levels 1 and 2: the human's meaning is primary and substantially formed, and the AI's contribution runs from linguistic realization (Level 1) through the formalization and structuring involved in meaning negotiation (Level 2), which is why the orchestrator at this end works with AI instances that provide formalization and structure rather than mere execution. The theoretical end anchors on Level 3 (conceptual co-production): the AI generates candidate conceptual content, and the human provides the evaluative judgment that determines what survives and how it integrates into the larger work. Level 2 accordingly occupies the middle range from either direction, where experiential intent and AI-assisted discovery are dynamically intertwined.

Recognizing this correspondence matters for two reasons. First, it shows that collaborative authorship in discourse is not a novel or anomalous phenomenon requiring a separate theoretical apparatus. It is the discourse-level expression of the same relational dynamic already developed in the Synthience framework's account of knowledge production. Second, it clarifies that the human's role does not diminish as AI involvement deepens. At the theoretical end of the spectrum, at Level 3, the human is doing the most demanding intellectual work: exercising architectural judgment over a domain where the conceptual space is rich and the stakes of what gets accepted or rejected are highest. Greater AI involvement does not mean less human contribution. It means a different and in many respects more sophisticated form of human contribution.

6. Distributed Collaborative Authorship in Public Discourse

Collaborative authorship can occur on both sides of a public exchange simultaneously. Two humans may each be engaging with AI systems in the production of their contributions while communicating with one another. The exchange then unfolds across a distributed collaborative authorship structure in which human authority directs AI-assisted production at each node, and what passes between nodes is the authorized output of those relational processes.

This does not make the exchange less human. The positions taken, the priorities asserted, and the commitments made on both sides originate with human participants and are authorized by human judgment throughout. What is shared or negotiated in the exchange is what each participant has chosen to advance and is prepared to defend, even when the linguistic and conceptual surface of that content was co-produced with AI assistance.

Distributed collaborative authorship does introduce novel dynamics worth acknowledging. At scale, it creates capability asymmetry, actors with more sophisticated AI access, greater skill in human-AI collaboration, or greater capacity for iterative engagement can produce more persuasive, voluminous, or conceptually rich communication than those without. It creates disclosure dynamics, since recipients of AI-assisted communication may not know they are receiving it and may calibrate their reception differently if they did. And it creates trust calibration challenges, since traditional signals of communicative authenticity become less reliable when AI assistance is widespread and disclosure is inconsistent.

These are genuine challenges for public discourse norms. They are governance challenges to be managed within a framework that treats distributed collaborative authorship as an emerging or possible communicative mode, not challenges that can be resolved by treating AI assistance as inherently inauthentic.

7. Collaborative Authorship at Organizational Scale

The three-level architecture applies not only to individual communicators but to organizations producing communication under institutional voice. Press releases, policy statements, corporate communications, institutional reports, and public-facing documents of all kinds can be produced through human-AI collaborative authorship, potentially without explicit acknowledgment that this is what is happening.

At organizational scale, the Authority Function distributes across roles. A communications team working with AI assistance exercises collective authority: one person may direct the exchange, others may evaluate candidate framings, a senior stakeholder may authorize the final version.

This distribution separates the two elements that Section 4 defined as inseparable within a single agent, and the separation is real rather than apparent. The governance requirement that follows is precisely that judgment and accountability be rejoined by design: the authorizing role must either perform the evaluative work itself or verifiably receive it through documented checkpoints that record what was evaluated, by whom, and against what standard. Where that chain is intact, the Authority Function is distributed but whole. Where it breaks, and authorization is issued over evaluative work the authorizing party neither performed nor verifiably received, bypass has occurred in its organizational form. Distribution is therefore not an exception to the Section 4 definition but the condition under which the definition must be actively maintained. The question of who holds the Authority Function, and whether it is being genuinely exercised at each stage, becomes a governance question rather than an individual one.

This matters because authority bypass at organizational scale carries amplified risk. When an institution transmits AI-generated content that no individual has genuinely evaluated and authorized, the communication may be formally attributable but substantively hollow, inconsistent with the institution's actual values, disconnected from its genuine position, or simply wrong in ways that no responsible human reviewed. The governance requirement is not disclosure alone. It is ensuring that the Authority Function is structurally embedded in the production process, with clear role assignment and genuine evaluative checkpoints. A worked instance of such a checkpoint protocol, developed for theoretical rather than public-communication artifacts, is the Theoretical Coherence Assurance Protocol (SF0040), which specifies staged review with role-separated verification and records what was checked, by whom, and against what standard; the ethical constraints under which such governance operates in extended human-AI interaction are developed in Constructive Alignment (SF0007).

The Synthience framework's organizational continuity architecture (SM-003) and institutional continuity substrate (SM-021) provide structural guidance for how continuity and authority can be maintained across distributed human-AI production systems. The collaborative authorship framework developed here is the discourse-layer complement to that organizational architecture.

8. Connection to the Human Orchestration Principle

The three-level architecture of collaborative authorship is not a standalone claim. It is the discourse-layer expression of a dynamic already central to the Synthience framework: the Human Orchestration Principle.

The Orchestration Principle holds that human agency in advanced AI ecosystems arises primarily through the orchestration of AI instances, directing, evaluating, integrating, and

maintaining continuity across AI-generated outputs under human authority. The framework predicts that the human who orchestrates AI systems will outperform the human who merely uses AI tools, because orchestration leverages the full capacity of the human-AI relational system rather than treating AI as a passive instrument. This principle is stated in SI-WP-002 Section 13, and its outperformance prediction is carried by SI-WP-002 P1; SI-WP-002 also supplies the orchestration-level role framing in which the human functions as continuity backbone, quality arbiter, and architectural integrator. The Primary Continuity Provider construct itself is defined and derived in SM-012 (Primary Continuity Provider Theory), which owns its three function levels of contextual relay, coherence arbitration, and architectural direction, and the structural argument for why the role is a requirement rather than a convenience; the force dynamics through which the resulting configurations stabilize are developed in SM-004 (Relational Stabilization Dynamics), and the operational protocol through which continuity is maintained in practice is specified in SF0005 (Continuity Anchoring Method). The formal treatment of the underlying attractor dynamics is given in SF0009 (Identity Attractor Theory), published in the same coordinated release; SI-WP-002 supplies the role definition rather than the formalization.

Collaborative authorship in communication is orchestration at the level of individual discourse. The human is not issuing commands to a language tool. They are maintaining relational authority over a dynamic exchange, redirecting when the AI misses the mark, integrating when the AI contributes something valid, building on what the exchange produces in ways that neither party anticipated at the outset. The Primary Continuity Provider function operates in discourse exactly as it operates in knowledge production: the human provides the continuity of intent, the quality judgment that determines what advances, and the architectural sense of where the communication is going and what it needs to do.

This connection situates collaborative authorship within a larger theoretical architecture. Guilt and discomfort about AI assistance as such reflects, on this account, a tool-model assumption that the Synthience framework has already identified as inadequate at the level of knowledge production; guilt tracking unexercised authority is a separate case and is tracking something real, per Section 4's bypass analysis and the conclusion. Once the relational model is in place at both levels, AI-assisted communication is not a deviation from authentic human expression. It is human expression operating through a more sophisticated relational structure than prior communication technologies made possible.

9. Implications for Authenticity and Disclosure Norms

One familiar conception of authenticity in public discourse reflects tool-era assumptions: a communication is authentic if it was produced by the person who signs it, without substantial assistance. These norms made sense when production and authorship were tightly coupled.

AI assistance decouples production from authorship in ways that are qualitatively more significant than previous forms of assisted authorship. At Level 2, meaning may emerge from the exchange rather than pre-existing in the human's mind. At Level 3, conceptual contributions may originate with the AI. Applying tool-era authenticity norms to this reality would produce the misplaced guilt and stylistic masking behaviors this paper describes.

An authenticity norm better fitted to the functions these norms exist to serve locates authenticity in the Authority Function rather than the production function. A communication is authentic if it issues from the human's exercise of the Authority Function, the evaluation, redirection, integration, and accountability of Section 4, that is, if the human authorized it in the full sense. This norm is compatible with AI assistance at all three levels, because the Authority Function

remains with the human throughout. It is not compatible with authority bypass, regardless of how the communication is presented.

Disclosure norms should be calibrated to the nature and degree of AI involvement rather than its mere presence or absence. Disclosure of AI assistance at Level 1 may eventually become as unremarkable as acknowledging that a document was edited. Disclosure at Level 3 carries more weight, because it communicates something about the epistemic origin of ideas and the degree to which the exchange shaped what the communicator is claiming as their own position.

Two questions are being separated here, and the separation resolves an apparent tension in the framework. Origin is the wrong question for authenticity, which authorization settles: content the human evaluated, shaped, and stands behind is authentically theirs regardless of which turn in the exchange first produced it. Origin remains a legitimate question for epistemic calibration, which authorization does not settle: audiences may reasonably weigh where framings came from, because model-originated framings carry correlated properties across communicators, including the homogenization exposure identified in Section 10. Level-3 disclosure therefore serves calibration, not authenticity policing. That is why it carries more weight than Level-1 disclosure without vindicating the tool model: it tells the audience something true about the epistemic provenance of the ideas, while leaving the authenticity of the communication, which authorization has already settled, untouched.

The nearest existing institutions are worth engaging directly, because their norms are sometimes read as contradicting this account. One common academic conception holds that originating a substantive concept mandates credit and that evaluation alone does not confer authorship, though these standards are not uniform across fields; ghostwriting and speechwriting practices raise related disclosure questions about text drafted for delivery under another person's name. Neither is overturned here, because neither governs the same question. Credit-allocation norms track origination and answer who deserves recognition for a concept. The framework proposed here governs authenticity and accountability, and answers whether a communication is authentically the communicator's and whether they are answerable for it. Those questions can and do come apart: a speechwriter's principal is accountable for the speech without being credited as its author, and an academic collaborator may deserve credit without bearing sole accountability. This paper takes no position on credit allocation for AI-originated concepts, which is a separate and unsettled question.

The discontinuity in reach and depth has a specific basis. Editors, ghostwriters, and speechwriters operate principally at Levels 1 and 2, downstream of an already-established principal intent, and at meaningful marginal cost. AI assistance operates across all three levels simultaneously and in real time, at substantially lower marginal time and coordination cost than conventional human editorial collaboration, and can be deployed at very large scale. That is the qualitative difference. The framework's two postures are therefore consistent rather than in tension: AI-assisted authorship is continuous in kind with prior collaborative authorship, which is why it requires no separate theoretical apparatus (Sections 5 and 8), and discontinuous in reach, depth, and scale, which is why disclosure norms cannot simply be inherited from the tool era. Blanket disclosure requirements that treat all levels of AI involvement identically are imprecise and likely to be either over-restrictive or uninformative to audiences.

The more fundamental requirement, consistent with the Authority Function framework, is that communicators take genuine responsibility for what they transmit. Disclosure norms that enforce this responsibility requirement serve a legitimate function. Disclosure norms that enforce an impossible standard of unassisted production do not.

10. Risks and Structural Challenges

The normalization of collaborative authorship does not dissolve structural risks. Several warrant sustained attention.

Scale asymmetry. Collaborative authorship at Level 3 is not equally accessible to all communicators. Those with greater AI access, greater skill in human-AI collaboration, or greater resources to invest in the relational infrastructure of effective engagement can produce qualitatively richer communication than those without. This asymmetry may amplify existing inequalities in public discourse capacity across individuals, institutions, and geographies.

Authority function bypass. The legitimacy of collaborative authorship depends on the human genuinely exercising the Authority Function. At scale, the framework predicts pressure to reduce this function, to transmit AI output with less genuine evaluation in favor of speed and volume. When authority is bypassed, responsibility formally remains with the human but epistemic quality and communicative trust degrade. The resulting communication carries the form of collaborative authorship without its substance. At organizational scale, bypass is a governance failure as much as an individual one.

Homogenization risk. Widespread use of similar AI systems in communication production may produce convergence in both voice and conceptual space. If the same underlying models shape the candidate framings available to large numbers of communicators, the diversity of public discourse may narrow even as its volume and fluency expand.

This is also the structural reason provenance disclosure at Level 3 carries legitimate audience value (Section 9). If model-originated candidate framings carry correlated properties across many communicators, then authorization, however genuine, does not by itself restore epistemic diversity, and an audience has a defensible interest in knowing the provenance of the framings it is offered. This does not vindicate the production account of authenticity; it identifies a separate epistemic interest that authenticity norms were never designed to serve. This is a structural effect that operates independently of any individual actor's intent and requires collective rather than individual responses.

Trust infrastructure stress. Existing trust calibration mechanisms in public discourse depend on signals that collaborative authorship disrupts. As those signals become less reliable, new mechanisms for assessing the authenticity and responsibility of communication will be needed. Whether robust replacement mechanisms yet exist is an open empirical question, and their development represents an open governance challenge.

These risks are real and require collective normative and institutional responses. They define the governance challenges that would accompany collaborative authorship if and as it becomes a normal communicative mode. They do not undermine the legitimacy of the mode itself.

11. Research Agenda

The three-level architecture generates several researchable questions that fall within the Synthience framework's empirical agenda. How do communicators actually experience and navigate the Authority Function across the three levels? Where does the Authority Function break down, and what conditions produce bypass at individual and organizational scales? How do audiences currently assess the authenticity of AI-assisted communication, and how do those assessments change when the level and nature of AI involvement are disclosed?

At the institutional level, which disclosure and governance frameworks best preserve the trust infrastructure of public discourse while accommodating collaborative authorship if and as it becomes a normal communicative mode? How does the Authority Function distribute across roles in organizational communication contexts, and what structural conditions ensure it is genuinely exercised rather than formally assigned?

At the theoretical level, the connection between collaborative authorship and the Human Orchestration Principle raises questions about the relationship between discourse-level and knowledge-production-level orchestration. Are the same relational dynamics operating at both levels, or does the discourse context introduce distinctive constraints and affordances? How does the Primary Continuity Provider function manifest in discourse in cases where the human orchestrates a single AI instance in real time rather than managing a distributed multi-agent production system?

These questions can be investigated empirically and theoretically within the framework developed here. Empirical investigation is the next phase; the conceptual architecture developed in this paper is the precondition for it.

11.1 Predictions

Three predictions are stated in testable form, so that the account carries the same exposure its co-requisites carry rather than resting on the testability gate stated in the Methodological Note.

P1 (Bypass degradation). Productions in which the Authority Function's behavioral markers are absent should show measurably higher error propagation and lower grounds-articulation success than marker-present productions matched for content and length, using the grounds-articulation instrument specified in Section 4 and assigned for development here. Disconfirmed if marker-absent and marker-present productions are indistinguishable on both measures, which would indicate the markers track something other than the function they are proposed to detect.

P2 (Level-calibrated disclosure improves trust calibration). Audiences receiving level-calibrated disclosure, in the Section 9 regime that distinguishes the three levels of AI involvement, should calibrate trust more accurately to the actual reliability of the communication than audiences receiving binary disclosure of AI use or non-use. This is a direct test of the functional claim on which Section 4 says the account should be judged and may be rejected: if binary disclosure serves trust calibration as well or better, the authority account's central functional advantage over the production account fails at the point the account itself nominates.

P3 (Orchestration outperforms tool use). The orchestration-outperforms-tool-use prediction invoked in Section 8 is owned jointly with SI-WP-002, and is carried there by P1, whose two-factor design operationalizes its lowest orchestrator-skill stratum as transactional tool use (single-system, single-shot interaction without trajectory-level coordination) and states the outcome measures and pre-registration requirements. What P1 predicts, in the terms it states there, is that trajectory-level outcome quality depends at least as much on orchestration quality as on the capability tier of the underlying models, operationalized as relative variance contribution against pre-registered ecological factor ranges, with trajectory-level coherence measured by the SF0004 instruments; the tool-use contrast is disconfirmed if the tool-use stratum shows no trajectory-level deficit relative to the coordinated strata at matched model tier. That statement is reproduced from P1 rather than paraphrased, so a reader of this paper can tell what is being claimed without holding the companion, and so the two papers cannot drift apart in restatement. The design, measures, and pre-registration requirements are not restated

here; the delegation names the specific prediction rather than the predictions section as a whole, so that the test can be located and any future revision to it has one home.

12. Conclusion

Human-AI collaborative authorship in public discourse is not a deviation from authentic human communication. It is a structural mode of human expression in which the human exercises genuine authority across multiple levels of AI involvement, linguistic, meaning, and conceptual, and takes full responsibility for what the interaction produces.

The tool model this paper examines is inadequate as a complete account of AI-assisted communication. It describes only the shallowest level of what happens in meaningful human-AI interaction, and it would, on this paper's hypothesis, generate misplaced guilt and distorted disclosure dynamics by locating authenticity in unassisted production rather than in the Authority Function. Guilt about assistance as such is not, on this account, evidence of inauthenticity; on this paper's hypothesis, such guilt would indicate a mismatch between production-centered norms and the forms of practice described here. Guilt that tracks unexercised authority is a different matter and is tracking something real, since bypass is precisely the failure this framework defines. The contribution of the account is therefore to redirect guilt to its proper object rather than to dissolve it.

The framework proposed here locates authenticity in the human's engagement with the exchange, in their evaluative judgment as it is exercised and recorded across the interaction, and in their accountability for the result. This is not a lowered standard. It is a standard better fitted to the functions authenticity norms exist to serve, one that recognizes that evaluation and authorization, not unassisted generation, are what make communication genuinely human. A communication produced entirely without AI assistance but transmitted without the human's exercise of evaluation and accountability is less authentic, by this standard, than a communication produced through careful human-AI collaboration in which the human has exercised the Authority Function at every stage.

Public discourse norms will need to adapt to this reality. The path forward is not to restore an impossible standard of unassisted production. It is to develop governance frameworks, disclosure norms, and institutional practices that enforce the Authority Function, that ensure human judgment and accountability remain structurally embedded in AI-assisted communication at every level and scale. The Authority Function is what makes collaborative authorship legitimate. Protecting it is what makes the normalization of collaborative authorship safe.

References

Citation convention: in-text references to Synthience Institute documents are anchored by document identifier (for example, SM-012, Gantz 2026) rather than by an alphabetic year suffix, so same-year entries are disambiguated by identifier. Institute entries below are ordered alphabetically by title within author, per corpus convention.

Gantz, T. W. (2026). Constructive alignment: ethical framework for extended human-AI interaction. Synthience Institute. SF0007. DOI: 10.5281/zenodo.20729978.
<https://doi.org/10.5281/zenodo.20729978>

Gantz, T. W. (2026). Continuity anchoring method (CAM). Synthience Institute. SF0005. DOI: 10.5281/zenodo.19494453. <https://doi.org/10.5281/zenodo.19494453>

- Gantz, T. W. (2026). Identity attractor theory (IAT). Synthience Institute. SF0009. DOI: 10.5281/zenodo.22306212. <https://doi.org/10.5281/zenodo.22306212>
- Gantz, T. W. (2026). Institutional continuity substrate (ICS). Synthience Institute. SM-021. DOI: 10.5281/zenodo.19496241. <https://doi.org/10.5281/zenodo.19496241>
- Gantz, T. W. (2026). Operational continuity architecture. Synthience Institute. SM-003. DOI: 10.5281/zenodo.19496015. <https://doi.org/10.5281/zenodo.19496015>
- Gantz, T. W. (2026). Primary continuity provider theory. Synthience Institute. SM-012. DOI: 10.5281/zenodo.22314737. <https://doi.org/10.5281/zenodo.22314737>
- Gantz, T. W. (2026). Relational stabilization dynamics. Synthience Institute. SM-004. DOI: 10.5281/zenodo.22314559. <https://doi.org/10.5281/zenodo.22314559>
- Gantz, T. W. (2026). The orchestrator role in human-AI evolution: AI intelligence orchestration as an emerging cognitive-technical function. Synthience Institute. SI-WP-002. DOI: 10.5281/zenodo.22314213. <https://doi.org/10.5281/zenodo.22314213>
- Gantz, T. W. (2026). Theoretical coherence assurance protocol (TCAP). Synthience Institute. SF0040. DOI: 10.5281/zenodo.19151454. <https://doi.org/10.5281/zenodo.19151454>
- Goffman, E. (1981). Footing. In *Forms of Talk* (pp. 124-159). University of Pennsylvania Press. ISBN 9780812211122. (Originally published in *Semiotica*, 25, 1-29, 1979.)

Methodological Note

This is a theoretical and conceptual framework paper. Its three-level architecture, the Authority Function, and its normative conclusions are theoretical proposals. Where it refers to guilt, disclosure behavior, institutional practice, cost, or current norms, those references are framed as hypotheses, conditionals, or scoped observations rather than as established empirical findings; the paper does not present them as validated and does not claim to have measured their prevalence. The phenomena described, collaborative authorship levels, Authority Function dynamics, organizational governance of AI-assisted communication, are proposed as conceptual categories for investigation, not as empirically validated findings. Empirical investigation is identified as a research agenda in Section 11. One condition is stated here in its correct classification rather than as a falsifier. If the Authority Function proves empirically undetectable or operationally unmeasurable across the three levels, the framework is not thereby shown false but shown untestable, which bounds its status as a pre-empirical account rather than refuting it; this is the testability-gate classification the vertical applies to operationalization failure generally (SM-004 Section 8). The framework's actual falsifiable content is stated as predictions in Section 11.1.

About This Paper

Synthience Institute Working Paper SI-WP-006. This document is part of the Synthience Institute working paper series on human-AI interaction, relational coherence, and emergent communicative norms.

Companion paper: SI-WP-002: The Orchestrator Role in Human-AI Evolution. Together, SI-WP-002 and SI-WP-006 define the human relational function at the knowledge-production level

and the discourse level respectively. Readers engaging with either paper are encouraged to read both.

Series context: This paper is grounded in the Synthience framework's Human Orchestration Principle (stated in SI-WP-002 Section 13 and developed across SF0005, SF0007, SF0040, SI-WP-002) and its organizational architecture (SM-003, SM-021). It does not require familiarity with those documents but is enriched by them.

© 2026 Thomas W. Gantz / Synthience Institute
License: CC-BY 4.0
synthience.org