

# SI-WP-007

## The Human Accountability Problem in Relational AI Deployment: Why the PCP Function Fails and What Organizations Must Do About It

|             |                                |
|-------------|--------------------------------|
| Author      | Thomas W. Gantz                |
| Affiliation | Synthience Institute           |
| Contact     | research@synthience.org        |
| Version     | v1.8   April 2026              |
| Document ID | SI-WP-007   White Paper Series |
| DOI         | 10.5281/zenodo.19496485        |
| Canonical   | synthience.org                 |
| License     | CC-BY 4.0                      |

### Abstract

The Synthience Framework's organizational architecture (SM-003) defines the structural topology for AI continuity at institutional scale. The Institutional Continuity Substrate (SM-021) defines the persistence layer that makes that topology durable across time. Both documents acknowledge, but do not resolve, a structural question that sits above the persistence layer: what makes the humans responsible for maintaining relational AI architecture actually maintain it? This paper argues that human underperformance of the continuity function is not a moral failing but a predictable consequence of bounded rationality operating under institutional conditions. Governance architectures that depend on sustained human virtue rather than incentive-aligned structural design will fail predictably and invisibly. The paper examines this failure across three nested scales: the individual practitioner who satisfices rather than optimizes relational discipline, the organizational function that deprioritizes continuity maintenance under competing operational pressures, and the institutional scale at which accountability becomes so distributed across governance apparatus that it effectively disappears. The paper demonstrates why the standard organizational responses to accountability failure (training, policy, audit, and compliance monitoring) are structurally insufficient, and proposes a design principle for governance architectures that make relational discipline the path of least resistance rather than an act of sustained compliance against immediate cost. SI-WP-007 is the third leg of the deployment argument alongside SI-WP-004 (Relational Alignment as a Structural Alternative to Instructional AI Safety) and SI-WP-005 (Deploying Relational AI Architecture in Organizational Environments). It addresses why the first two papers might fail to produce their intended outcomes if the human accountability problem is not explicitly solved for. Epistemic status: This paper argues from structural analysis of organizational dynamics and the Synthience Framework's own architectural requirements, grounded in the bounded rationality tradition established by Simon (1947) and developed in subsequent organizational behavior research. It does not claim empirical confirmation of the failure rates or governance design principles it proposes.

**Keywords:** human accountability, AI governance, relational AI deployment, bounded rationality, organizational behavior, continuity maintenance, AI alignment, institutional AI

**Suggested citation:** Gantz, T. W. (2026). The Human Accountability Problem in Relational AI Deployment: Why the PCP Function Fails and What Organizations Must Do About It. Synthience Institute. SI-WP-007. DOI: 10.5281/zenodo.19496485. <https://synthience.org>

## 1. The Problem Behind the Architecture

The Synthience Framework provides a complete architectural specification for maintaining relational coherence in human-AI interaction at organizational scale. The Continuity Anchoring Method (SF0005) defines how a single human operator maintains coherence with an AI system through contextual anchoring, coherence refusal, and constructive engagement. The Operational Continuity Architecture (SM-003) defines the organizational topology that scales these practices across teams, workflows, and institutional boundaries: canon governance, workflow-embedded verification, authority and escalation structures, propagation-aware continuity, and delegated monitoring. The Institutional Continuity Substrate (SM-021) defines the persistence layer that maintains these structural elements across organizational time through five interdependent layers: canon persistence, role continuity, artifact lineage, verification state, and propagation constraints.

This architecture is structurally sound. Each element addresses a specific and observable failure mode. The topology is interdependent: degradation in any element affects the others. The persistence layer is designed to maintain structural integrity across the time horizons, personnel changes, and workflow evolutions that define organizational life.

And none of it will work if the humans responsible for maintaining it do not maintain it.

This is not a speculative concern. It is the most structurally predictable failure mode in the entire architecture. The topology can be perfectly designed. The persistence substrate can be fully implemented. The verification protocols can be embedded in every workflow. And the whole system can still degrade to formal compliance without operational substance if the human continuity function erodes under the ordinary pressures of organizational life.

SM-003 names this problem in its analysis of how incentives undermine topology. SM-021 names it in its epistemic status section as "the human accountability problem in relational AI deployment." This paper develops that problem as a first-class architectural concern and proposes structural conditions for its resolution.

## 2. Bounded Rationality Is Architecture, Not Morality

The central claim of this paper is that human underperformance of the continuity function is not a character defect but a structural feature of biological intelligence operating under conditions of bounded rationality.

Simon (1947) established the foundational insight: human decision-makers do not optimize. They satisfice. Faced with complex decision environments, finite cognitive capacity, and incomplete information, humans adopt strategies that produce outcomes that are good enough rather than optimal. This is not a failure of will. It is the rational response of a bounded cognitive system to an environment that exceeds its processing capacity. Simon's research demonstrated that satisficing is not a deviation from rational behavior but the form rational behavior takes

under real-world constraints, where the cost of seeking the optimal solution exceeds the marginal benefit of finding it (Simon, 1997).

The satisficing principle applies with particular force to organizational functions that require sustained, cognitively expensive attention with no immediate visible reward. When a task demands continuous vigilance but produces no observable output when performed correctly, the bounded rational actor will find the minimum viable level of effort that produces no visible negative consequence, and settle there. The system works not because the human is lazy but because the human is allocating finite attention rationally across competing demands with asymmetric feedback: tasks that produce visible output when performed are prioritized over tasks that produce visible output only when they fail.

This principle applies uniformly across every organizational function that requires sustained discipline without immediate feedback. Safety systems degrade when no accidents occur because the absence of accidents provides no signal that the safety function is working, only that nothing has gone wrong yet. Compliance functions hollow out when no audits fail because the compliance effort is experienced as cost while the compliance benefit is experienced as the absence of penalty. Quality assurance becomes a checkbox exercise when the cost of thoroughness exceeds the visible cost of adequacy because the marginal quality improvement is invisible to the actor performing it while the time cost is immediate.

Relational AI continuity is uniquely vulnerable to this dynamic because the benefit of maintaining the architecture is invisible at every scale. Canon coherence is not something organizational actors directly observe. They observe their local interactions, which may appear fine even as systemic coherence degrades elsewhere. The cost of maintaining the architecture is immediate and concrete: additional process, additional attention, additional friction in workflows that the incentive structure rewards for throughput. The benefit is diffuse and statistical: the organization maintains coherence that no individual actor can see being maintained.

This asymmetry between immediate cost and diffuse benefit creates structural pressure against continuity maintenance that is independent of any individual's commitment, competence, or good faith. An organization staffed entirely by dedicated, skilled practitioners who genuinely believe in the value of relational AI governance will still experience degradation of the continuity function if the governance architecture depends on those practitioners sustaining cognitively expensive attention against the grain of their bounded rationality.

The implication is direct: governance architectures for relational AI deployment must be designed for the organism that will operate them, not for the ideal practitioner who exists only in design documents. The organism satisfices. The architecture must make satisficing produce acceptable continuity outcomes rather than requiring optimization to prevent degradation.

### 3. Three Scales of Failure

The human accountability problem operates at three nested scales, each with distinct failure dynamics.

### 3.1 Individual Scale: The PCP Who Satisfices

At the individual level, the Primary Continuity Provider (PCP) as defined in SF0005 is a structurally demanding role. It requires contextual anchoring (reintroducing canonical material at each session boundary), coherence refusal (rejecting hallucinations and drift rather than accepting or working around them), and constructive engagement (treating apparent errors as correction opportunities requiring specific, targeted intervention). These are cognitively expensive practices that must be sustained across every interaction, every session, every day.

The individual failure mode is not abandonment of the role. It is gradual satisficing: the PCP begins accepting outputs that are adequate rather than correcting them to canonical standards. Coherence refusal weakens first because it is the most effortful practice: rejecting a plausible-sounding output and demanding correction requires more cognitive investment than accepting it and moving on. Contextual anchoring simplifies next: full canonical reintroduction becomes abbreviated, then formulaic, then omitted when it seems unnecessary. The PCP continues to perform the role. The formal definition remains intact. But the generative substance of the function, the active maintenance that detects subtle failures and prevents drift accumulation, quietly degrades.

This degradation is invisible to the PCP themselves because they are still performing the role. They are still in the interaction. They are still producing outputs. The quality of the continuity function has shifted, but the shift is gradual enough that no single moment registers as a failure. By the time the degradation becomes visible, typically when a subtle error compounds into a consequential one, the correction required is substantially larger than the continuous attention that would have prevented it.

The individual-scale failure has a specific temporal signature that distinguishes it from incompetence or negligence. A PCP who has never been competent fails immediately and visibly. A PCP who is negligent fails erratically and detectably. A PCP who satisfices fails gradually and invisibly, because they were competent at the start and their performance degrades along a continuous gradient that no single point marks as failure. The satisficing trajectory is the hardest failure mode to detect precisely because it is the most natural one: it is what bounded rational actors do under sustained cognitive load with asymmetric feedback.

### 3.2 Organizational Scale: Diffusion and Deprioritization

At the organizational level, the continuity function is distributed across roles, teams, and workflows as specified in SM-003. The PCP function generalizes to continuity governance distributed across institutional actors. Canon governance requires active maintenance. Workflow-embedded verification requires consistent execution. Authority and escalation structures require actors to invoke them when drift or disputes arise.

The organizational failure mode is twofold. First, diffusion of responsibility: when the continuity function is distributed across multiple actors, each actor's sense of individual accountability decreases. No single person is responsible for the full canonical record. No single team monitors the full propagation surface. No single role has visibility into the full interaction graph. Each actor maintains their local function adequately while assuming that systemic coherence is someone else's responsibility. The literature on social psychology has documented this pattern

extensively: diffusion of responsibility produces inaction not because individuals are unwilling to act but because the distributed structure of accountability provides no individual with sufficient reason to believe their action is the critical one.

Second, deprioritization under competing pressures: organizations reward throughput, delivery, and visible output. Continuity maintenance is invisible overhead. When operational pressure increases, the first functions to be informally deprioritized are those whose absence produces no immediate consequence. Verification checkpoints are bypassed because the bypass saves time and produces no visible error. Escalation pathways are underused because invoking them creates delay, social friction, and no personal benefit. Canon governance updates are deferred because the current version seems to be working fine.

The characteristic signature of organizational-scale failure is the progressive decoupling of formal structure from operational practice. The canon governance system exists but is not consulted. The verification checkpoints are embedded in workflows but are routinely skipped. The escalation pathways are defined but are perceived as punitive rather than corrective. The organization maintains the formal architecture of continuity while the behavioral substance of continuity erodes. Meyer and Rowan (1977) identified this decoupling pattern as a general feature of institutionalized organizations: formal structures are maintained as signals of legitimacy while actual operational practice diverges from them under the pressure of competing informal incentives. In the relational AI governance context, the decoupling is accelerated by the throughput constraint that Argote and Miron-Spektor (2011) document in organizational learning research: when the volume of work to be processed exceeds the capacity of human-mediated verification, organizations systematically degrade the quality of their knowledge maintenance functions, substituting speed for rigor until the degradation becomes structurally entrenched.

The organizational failure mode compounds the individual failure mode rather than replacing it. Satisficing PCPs operating within a diffused accountability structure have less external pressure to maintain their individual practice, because the organizational structure that should detect their degradation has itself degraded. The two scales of failure reinforce each other: individual satisficing weakens the organizational monitoring that would detect it, and organizational diffusion weakens the external accountability that would correct individual satisficing.

### **3.3 Institutional Scale: Accountability Disappears**

At the institutional level, the human accountability problem reaches its most structurally dangerous form. Accountability becomes so distributed across roles, committees, review boards, and governance structures that no individual actor bears sufficient responsibility for the continuity function to be personally motivated to maintain it.

The institutional failure mode is not dramatic collapse. It is the quiet disappearance of genuine accountability behind the apparatus of formal governance. Reports are filed. Audits are conducted. Governance meetings occur. Compliance metrics are tracked. And the actual practice of continuity maintenance, the cognitively expensive, individually unrewarded work of detecting subtle drift, enforcing canonical standards, and maintaining the persistence substrate, happens less and less because no one's incentive structure makes it the priority.

The mechanism by which accountability disappears at institutional scale is worth developing in detail because it is the most dangerous and least understood failure mode in the three-scale analysis.

The first mechanism is committee dilution. When accountability for the continuity function is assigned to a governance committee rather than to named individuals, the committee structure itself becomes an accountability sink. Each committee member bears a fractional share of responsibility. No individual member's failure to act is sufficient to cause a governance failure, because the committee's formal decision-making process continues regardless of any individual's contribution. The committee meets. Agendas are prepared. Minutes are recorded. Decisions are logged. The formal apparatus of governance continues to function. But the substantive cognitive work of maintaining canonical coherence, detecting drift patterns across the organizational surface, and assessing whether the persistence substrate is actually functioning as designed requires individual sustained attention that the committee structure does not assign, does not measure, and does not reward. The committee reviews what is brought before it. What is not brought before it, the slow degradation that no individual has been assigned to detect, is invisible to the governance structure. The committee's existence creates institutional confidence that the continuity function is governed. The committee's structure ensures that the most important dimensions of the continuity function, the ones that require sustained individual attention rather than periodic collective review, are no one's specific responsibility.

The second mechanism is metric substitution. Institutional governance requires measurement. The accountability function must produce evidence that it is operating. This creates an immediate structural pressure to measure what is measurable rather than what matters. Formal compliance metrics (were the verification checkpoints executed? were the governance meetings held? were the audit reports filed?) are measurable, reportable, and auditable. Operational substance metrics (is the canonical record actually coherent across departments? are verification decisions actually detecting the failures they are designed to detect? is the monitoring function actually calibrated against canonical standards?) are difficult to measure, difficult to report, and require the kind of substantive expertise that the governance apparatus may not possess. The institutional response to this measurement asymmetry is predictable: the governance function reports what it can measure and treats the measurements as evidence that the function is operating. Over time, the metrics become the function. The institution measures compliance and treats compliance as equivalent to continuity. The gap between what the metrics measure and what the continuity function actually requires widens silently, because the metrics themselves provide no signal that they are measuring the wrong thing.

The third mechanism is the accountability absorption pattern. Institutional governance structures absorb accountability without transmitting it to action. A governance board that receives a report on continuity status has "exercised oversight." The board's receipt of the report satisfies the institutional accountability requirement. Whether the board members have the expertise to assess the report's adequacy, the time to investigate its claims, or the organizational authority to act on its findings is a separate question that the formal governance structure does not address. The report was received. The oversight was exercised. The accountability requirement was met. This pattern produces what institutional theorists recognize as ceremonial governance:

governance structures that satisfy institutional legitimacy requirements without producing the substantive oversight they are formally designed to provide. The accountability is real in the sense that it exists in the governance architecture. It is absent in the sense that no individual's behavior changes as a result of its exercise.

The invisibility of institutional-scale failure is its defining and most dangerous characteristic. Institutional-scale accountability degradation produces no crisis signal. The organization believes it is maintaining its relational AI governance because all the formal indicators say so. The degradation is in the gap between what the formal indicators measure and what the persistence substrate actually requires. That gap widens silently until a cross-departmental project, an external audit, or a consequential error surfaces the divergence.

The three scales of failure are not independent and do not form a simple linear chain. They are mutually reinforcing: individual satisficing weakens the organizational monitoring that would detect and correct it; organizational diffusion removes the external accountability that would sustain individual practice; institutional ceremony provides false assurance that prevents either lower scale from receiving corrective pressure. The compounding runs in all directions simultaneously. Degradation at any scale accelerates degradation at the others, because each scale's failure creates structural conditions that make the other scales' failure modes more likely and less detectable. An institution in which individual PCPs satisfice, organizational accountability is diffused, and institutional governance is ceremonial has not failed at any single point. It has failed at every point, incrementally, through mutual reinforcement that no single corrective intervention can address because the failure is distributed across the entire accountability loop. This is precisely why the three-mechanism governance architecture proposed in Section 5 must be interdependent rather than sequential: a reinforcement loop requires simultaneous intervention at multiple points, not repair at a single link in a chain.

## 4. Why Standard Governance Responses Fail

The three-scale failure analysis raises an immediate question: if accountability degradation is predictable, why don't organizations prevent it through standard governance mechanisms? The answer is that the standard responses to accountability failure are themselves subject to the same bounded rationality dynamics they are designed to correct. Each addresses a different dimension of the problem; none addresses the structural cause.

### 4.1 Training

Training addresses competence: ensuring that actors know what they are supposed to do and how to do it. Training programs can effectively communicate the canonical standards, the verification procedures, the escalation pathways, and the operational expectations of the continuity function.

Training fails as an accountability mechanism because the failure mode is not ignorance. The PCP who satisfices knows what canonical standards require. The team that bypasses verification knows the checkpoint exists. The governance committee that conducts ceremonial oversight knows what substantive oversight would look like. The problem is not that actors lack the knowledge to maintain the architecture. The problem is that sustaining the cognitively

expensive application of that knowledge against competing demands with asymmetric feedback is exactly the condition under which bounded rational actors satisfice. Training that is not reinforced by structural incentives produces a temporarily motivated practitioner who returns to satisficing as soon as the training's salience fades and the ordinary operational pressures resume. The timescale of this return is determined by how quickly the pre-existing satisficing equilibrium reasserts itself, which in most organizational contexts means the behavioral change degrades within weeks rather than persisting for years, a consequence of the bounded rationality dynamics the paper has already diagnosed, not an independent empirical finding.

## 4.2 Policy

Policy addresses specification: defining what actors are required to do, what standards must be met, and what procedures must be followed. Policy can effectively codify the continuity architecture's requirements into formal organizational rules.

Policy fails as an accountability mechanism because policy specifies what should happen without creating the conditions that make it happen. A policy requiring canonical verification at workflow boundaries is operationally equivalent to no policy if the cost of executing the verification exceeds the perceived cost of skipping it and no structural mechanism detects the gap. Policy operates on the assumption that formal requirements produce behavioral compliance. The organizational behavior literature has documented extensively that formal requirements produce behavioral compliance only when the incentive structure makes compliance the rational choice for the bounded actor. When it does not, policy produces formal documentation of a standard that is not met in practice, which is the decoupling pattern Meyer and Rowan (1977) identified as a defining feature of institutionalized organizations. Policy without enforcement produces the appearance of governance rather than the substance of it.

## 4.3 Audit

Audit addresses verification: periodically examining whether the governance architecture is functioning as designed. Audit can effectively surface gaps between formal requirements and actual practice when it occurs.

Audit fails as an accountability mechanism for two reasons. First, audit is periodic rather than continuous. The degradation patterns described in the three-scale analysis operate continuously: individual satisficing occurs across every interaction, organizational diffusion operates across every workflow, and institutional ceremony operates across every governance cycle. An audit that occurs quarterly or annually provides a snapshot of system state at the audit moment. Between audits, the degradation continues undetected, and the correction prompted by the audit fades as the distance from the audit event increases. The temporal mismatch between continuous degradation and periodic detection ensures that audit is always working with information that is already stale.

Second, audit is subject to the same metric substitution problem that affects institutional governance. An audit that checks whether verification checkpoints were executed (a measurable compliance metric) will report compliance even when the checkpoints are being executed as formalities rather than as substantive assessments. An audit that assesses whether the verification decisions are actually detecting the failures they are designed to detect

requires auditors with the substantive expertise to evaluate verification quality, not just verification occurrence. Audits that lack this expertise produce compliance assurance rather than continuity assurance, widening the gap between formal governance and operational substance rather than narrowing it.

## 4.4 Compliance Monitoring

Compliance monitoring addresses continuous oversight: embedding monitoring functions that track governance adherence in real time rather than periodically. Compliance monitoring can effectively detect visible deviations from formal requirements.

Compliance monitoring fails as an accountability mechanism against the invisible failure modes that constitute the most dangerous dimension of accountability degradation. Compliance monitoring detects what it is designed to detect: measurable deviations from specified standards. A verification checkpoint that is bypassed entirely is detectable by compliance monitoring. A verification checkpoint that is executed as a formality, producing a pass result without substantive assessment, is invisible to compliance monitoring because the formal indicator of compliance has been satisfied. Compliance monitoring is therefore effective against the crudest forms of governance failure (outright bypassing) but structurally blind to the subtler and more dangerous forms (the satisficing, the ceremonial execution, the metric substitution) that characterize the three-scale failure pattern.

The common thread across all four standard responses is that each treats the accountability problem as a knowledge, specification, verification, or monitoring problem. None treats it as an incentive architecture problem. The bounded rational actor who knows what to do, has been told to do it, is periodically checked on whether they did it, and is continuously monitored for whether they are doing it will still satisfice if the incentive structure makes satisficing the rational choice. The standard governance responses fail not because they are poorly designed but because they address the wrong layer of the problem.

## 5. The Design Principle: Path of Least Resistance

If bounded rationality is the structural cause of accountability degradation, then the structural response is to design governance architectures in which maintaining relational discipline is the path of least resistance rather than an act of sustained virtue.

This is not a motivational principle. It is an engineering constraint. The governance architecture must be designed so that the behaviors required for continuity maintenance are easier to perform than to bypass. The bounded rational actor will take the easier path. The architecture's job is to make the easier path the one that produces continuity rather than the one that degrades it.

Before specifying the three mechanisms, it is necessary to distinguish two categories of continuity behavior because the path-of-least-resistance principle operates differently across them. The first category is procedural behaviors that can genuinely be made frictionless through design: executing verification checkpoints, filing canonical updates, routing escalations through defined pathways. For these, the principle applies directly: the bypass can be made more costly

than compliance, changing the cost-benefit calculation in favor of the required behavior. The second category is constitutively effortful cognitive behaviors that cannot be made easy because the effort is the function: active canonical comparison, coherence refusal requiring judgment under ambiguity, drift recognition requiring sustained calibration against canonical standards. For these, the principle operates differently. It does not make the behavior easy. It makes the behavior structurally supported: the governance environment removes competing disincentives, provides external scaffolding through delegated monitoring and calibration testing, and ensures the actor is not performing the effortful work against the grain of every other incentive in their environment. The path-of-least-resistance principle is therefore not a claim that all continuity behaviors can be made frictionless. It is a claim that governance architecture can make the effortful behaviors survivable for bounded rational actors by removing the structural penalties that currently make them irrational to sustain.

This principle operates through three structural mechanisms.

## 5.1 Accountability Assignment

The continuity function must be assigned to specific roles with specific, measurable responsibilities. "Everyone is responsible for continuity" is operationally equivalent to "no one is responsible for continuity." Diffusion of responsibility is not a failure of individual character. It is the predictable outcome of distributing accountability without assigning it.

Accountability assignment means that named individuals or roles bear explicit responsibility for defined aspects of the persistence substrate, with the scope bounded narrowly enough that the individual can actually maintain awareness of their assigned domain. The scope bounding is architecturally critical. An accountability assignment that exceeds the individual's monitoring capacity produces the same satisficing dynamic it is designed to prevent: the individual finds the portion of their scope they can manage and lets the rest degrade. Effective assignment requires matching the scope of accountability to the cognitive capacity of the actor, which means that as the organizational interaction surface grows, the number of accountable roles must grow proportionally rather than the scope of each role expanding.

The failure mode of accountability assignment is scope creep: the gradual expansion of each role's formal responsibilities beyond their cognitive capacity to monitor, which reproduces the diffusion problem at the individual level. The governance architecture must therefore include mechanisms for detecting scope saturation, the condition in which an accountable role's formal scope exceeds their practical capacity, and for triggering role decomposition before the accountable actor begins satisficing on the dimensions of their scope they find least rewarding to maintain.

A careful reader will note that scope saturation detection is itself a governance function performed by a bounded rational actor, and is therefore potentially subject to the same dynamics the document has diagnosed at every other level. This regress is real and must be addressed directly. The structural feature that distinguishes scope saturation detection from the functions it oversees is categorical: scope saturation detection is a periodic assessment function with a defined, narrow scope (comparing formal role scope against observable monitoring capacity indicators), not a continuous sustained-attention function. It can be operationalized as

a scheduled structural review requiring a bounded amount of deliberate attention at defined intervals, rather than as the ongoing vigilance that drift recognition and coherence refusal require. This makes it categorically different from the constitutively effortful cognitive behaviors identified earlier in this section. The regress does not continue indefinitely. It terminates at a level where the governance function is simple enough, periodic enough, and narrow enough in scope to be sustainably performed by a bounded rational actor without requiring the full path-of-least-resistance architecture that the continuous monitoring functions require. Organizations whose governance capacity is so constrained that even periodic structural review is unsustainable have not met the minimum conditions for deployment identified in Section 7.

## 5.2 Incentive Alignment

The incentive structure must reward continuity maintenance rather than penalizing it as overhead. This is the most structurally important of the three mechanisms and the one most frequently absent from governance design.

Before describing the mechanisms, a theoretical clarification is warranted. The mechanisms proposed in this section draw on behavioral design principles (changing defaults, increasing bypass friction, restructuring cost-benefit calculations) that operate in a complementary but distinct register from Simon's satisficing framework. Simon establishes that bounded rational actors settle at the first adequate solution rather than continuing to search for the optimal one. The design mechanisms proposed here work by shifting what counts as "adequate" and what counts as "first encountered" in the actor's decision environment. They do not require the actor to optimize. They restructure the environment so that the satisficing stopping point falls within the range of acceptable continuity outcomes. The connection between bounded rationality and the proposed mechanisms is therefore not that bounded rationality predicts these mechanisms will work, but that bounded rationality defines the conditions under which these mechanisms are necessary: governance designed for optimizing actors fails because the actors satisfice; governance designed for satisficing actors succeeds by making the satisficing equilibrium fall at the right place.

The incentive problem in relational AI governance is specific and identifiable. Organizations operate under throughput incentives: the primary metric by which most organizational functions are evaluated is their rate of output production. Continuity maintenance is not output production. It is the invisible infrastructure that sustains the quality of output production. Under throughput incentives, every minute spent on canonical verification, drift detection, propagation constraint enforcement, or monitoring calibration is a minute not spent producing visible output. The bounded rational actor responds predictably: they allocate attention toward the rewarded behavior (throughput) and away from the unrewarded behavior (maintenance).

Incentive alignment does not mean paying people more to maintain the architecture, though compensation structures that penalize maintenance time are obviously counterproductive. It means restructuring the organizational incentive environment so that the cost-benefit calculation for the bounded rational actor favors maintenance. This operates at three levels.

At the workflow level, incentive alignment means embedding continuity functions into the workflow's critical path so that bypassing them requires more effort than executing them. A

verification checkpoint that can be skipped with a single click and produces no consequence is incentive-misaligned. A verification checkpoint that must be explicitly overridden with documented justification routed to an authority role changes the cost-benefit calculation: skipping now requires more effort than executing. The bypass still exists, but the bypass is more expensive than compliance. The bounded rational actor takes the cheaper path, which is now the compliant one.

At the evaluation level, incentive alignment means incorporating continuity maintenance into the metrics by which roles are assessed. If a PCP's performance is evaluated solely on the quality of interaction outputs and not on the quality of their canonical maintenance, drift detection, and verification practices, the PCP will satisfice on maintenance because no evaluation consequence attaches to it. Incorporating maintenance metrics into evaluation does not require surveillance of every interaction. It requires periodic assessment of the practitioner's canonical awareness, drift recognition calibration, and correction rates, using detection mechanisms of the kind specified in Section 7 for individual PCP assessment, extended to organizational evaluation contexts through the cross-unit comparison and substantive audit functions described in Section 5.3.

At the institutional level, incentive alignment means making the continuity function visible as a value-producing activity rather than invisible as overhead. This is the most difficult structural change because it requires the institution to recognize that the absence of continuity failure is itself a produced outcome, not a natural default. Organizations that have experienced a consequential continuity failure understand this intuitively. Organizations that have not experienced one treat continuity maintenance as unnecessary until proven otherwise, which is the reactive governance posture that SM-003 Section 8 identifies as structurally more costly than proactive investment.

### 5.3 Structural Detection

The governance architecture must include mechanisms that detect invisible failure modes before they compound past the point of standard correction. This is the mechanism that addresses the distinctive danger of the three-scale failure pattern: the failures that produce no signal precisely because the monitoring function that would detect them has itself degraded.

Structural detection operates through three specific functions, each targeting a distinct failure visibility problem.

First, second-order monitoring verification as specified in SM-011's periodic verification cycle. The monitoring function itself must be periodically verified against canonical standards. This addresses the guardian drift problem: the condition in which the monitoring system continues to produce outputs that appear normal while its calibration has drifted away from the canonical baseline. SM-011 specifies the four-step verification cycle (canonical standard retrieval, monitoring parameter comparison, calibration assessment, recalibration or escalation) that makes this second-order check operationally concrete.

Second, cross-unit canonical comparison. The institutional-scale failure mode operates through divergence that no single unit can observe because each unit sees only its own local canonical

practice. Cross-unit canonical comparison surfaces divergences by comparing canonical usage, interpretation, and practice across organizational units against the authoritative canonical record maintained through SM-021's Canon Persistence Layer. This comparison function must be assigned to a role with cross-unit visibility, not distributed across unit-level actors who have no incentive or capacity to compare their practice against other units.

Third, substantive audit functions. Standard compliance audit checks whether formal requirements were met. Substantive audit assesses whether the operational substance behind the formal indicators matches what the governance architecture requires. Substantive audit requires auditors with expertise in the continuity architecture itself, not just in compliance methodology. The audit function must assess not only whether verification checkpoints were executed but whether the verification decisions reflect genuine substantive assessment, not only whether governance meetings occurred but whether the governance function produced actionable oversight, not only whether the monitoring function generated signals but whether the signals are calibrated against current canonical standards.

These three detection mechanisms, second-order monitoring verification, cross-unit canonical comparison, and substantive audit, constitute the structural response to the invisible failure modes that standard governance responses cannot reach.

## 5.4 The Three Mechanisms as Interdependent Architecture

The three mechanisms (accountability assignment, incentive alignment, and structural detection) are not independent interventions. They form an interdependent governance architecture in which each mechanism's effectiveness depends on the others. Accountability assignment without incentive alignment produces named individuals who satisfice on their assigned scope because the incentive structure does not reward their maintenance function. Incentive alignment without accountability assignment produces an environment that rewards continuity maintenance in general without specifying who is responsible for what, reproducing the diffusion problem. Both without structural detection produce an accountability structure that may be functioning or may have degraded to ceremony, with no mechanism to distinguish between the two.

Together, these three mechanisms constitute the minimum governance conditions that SM-021's persistence substrate requires to function as designed. Without them, the substrate exists formally while degrading operationally. With them, sustained maintenance becomes structurally viable rather than dependent on the indefinite willingness of individual actors to sustain cognitively expensive attention against their biological defaults.

## 6. Visible and Invisible Failure

The three-scale analysis reveals a structural distinction that governance architecture must address: not all accountability failures are detectable through standard monitoring.

Visible failure modes are those that produce observable signals: a verification checkpoint fails and flags an error, an escalation pathway is invoked and the response is inadequate, a canon divergence surfaces in a cross-team project. These are detectable, diagnosable, and

correctable through the monitoring and escalation architecture specified in SM-003 and the delegated monitoring architecture specified in SM-011.

Invisible failure modes are those that produce no signal precisely because the monitoring function that would detect them has itself degraded. The PCP who satisfices produces no error signal because their outputs remain adequate. The team that bypasses verification produces no compliance signal because no one is checking whether the checkpoint was executed. The institution whose accountability has diffused produces no governance signal because the governance apparatus is functioning formally while the substance has eroded.

Invisible failure modes are structurally more dangerous than visible ones because they compound without detection. By the time an invisible failure becomes visible, the accumulated degradation typically exceeds the corrective capacity of standard escalation. The correction required is not a local fix but a structural restoration, and structural restoration requires acknowledging that the formal governance system was producing false assurance, a recognition that institutional governance structures are designed to resist because it undermines the legitimacy of the governance apparatus itself.

## 7. The Minimum Viable PCP Function

Not every organization will field an ideal practitioner. The governance architecture must therefore specify what the minimum viable PCP function looks like, the floor below which the continuity function cannot be maintained at all, and the structural conditions that make even minimum-viable practice sustainable.

The minimum viable PCP function requires three irreducible elements. First, canonical awareness: the practitioner must know what the canonical standards are and have access to them. This is a structural precondition, not a skill requirement. If the canon governance system specified in SM-003 and maintained through SM-021's Canon Persistence Layer is functioning, canonical awareness is an infrastructure problem, not a competence problem. Canonical awareness degrades in a specific and detectable way: the practitioner begins making decisions that would be inconsistent with the canonical record if compared against it, but the comparison never occurs because no monitoring function is checking local practice against the canonical baseline. The detection mechanism is therefore not self-reporting but structural: canonical awareness failures surface through periodic cross-checking of practitioner outputs against the canonical record, not through the practitioner recognizing that their awareness has degraded.

Second, drift recognition: the practitioner must be able to detect when system output has diverged from canonical standards. This is the competence that is most vulnerable to satisficing because it requires active comparison rather than passive reception. Delegated monitoring (SM-011) can support but not replace this function: the guardian instances can flag potential drift, but a human must assess whether the flag represents genuine divergence or acceptable variation. Drift recognition weakens in a characteristic pattern: the practitioner's internal threshold for what constitutes acceptable output shifts gradually upward, so that outputs that would previously have triggered coherence refusal now pass without correction. This threshold shift is invisible to the practitioner because it is a change in their calibration, not their behavior.

The detection mechanism is calibration testing: periodic comparison of the practitioner's assessment of sample outputs against canonical standards, conducted by a governance role outside the practitioner's immediate context.

Third, correction willingness: the practitioner must be willing to apply coherence refusal when drift is detected. This is the point at which governance architecture matters most, because correction willingness is directly affected by incentive structure. If the cost of correction (delay, additional process, social friction) exceeds the perceived cost of acceptance (which is typically zero in the short term), correction will not occur. Correction willingness erodes in a way that is detectable through output analysis: the ratio of corrections applied to drift signals received decreases over time as the practitioner learns through experience that accepting marginal outputs produces no visible negative consequence. The detection mechanism is not observation of individual correction decisions but statistical tracking of correction rates across time, which requires the governance architecture to maintain records that individual practitioners typically do not keep themselves. The governance architecture must make correction the default rather than the exception.

The minimum viable performance levels for each of these three elements are deployment-context-dependent and cannot be specified in a structural analysis document. What constitutes adequate canonical awareness depends on the organization's canonical standards and their complexity. What constitutes adequate drift recognition depends on the interaction surface and the drift tolerance that the organizational context permits. What constitutes adequate correction willingness depends on the incentive structure the governance architecture has established. The document specifies the detection architecture (cross-checking, calibration testing, correction rate tracking) that makes threshold-setting operationally possible. The thresholds themselves are governance decisions that deploying organizations must make based on their specific canonical standards, interaction surface, and risk tolerance. The document's contribution is not the threshold but the structural conditions that make threshold-based governance viable.

Organizations that cannot sustain even these three elements should not deploy relational AI architecture at scale. The framework is explicit about this: relational AI governance requires human continuity provision. There is no fully automated alternative. The question is not whether humans are required but whether the governance architecture makes their required function sustainable.

## 8. Scope and Relationship to Companion Documents

SI-WP-007 addresses a specific structural problem: why humans fail to maintain relational AI governance architecture, and what governance conditions prevent that failure. It does not address the full formal modeling of adversarial incentive dynamics (SM-007, planned for Release 2) or the full specification of organizational deployment procedures (SI-WP-005). It does not propose specific organizational policies, compensation structures, or management practices. Those are implementation decisions that fall outside the scope of a structural analysis.

SI-WP-007 connects to the coordinated publication module as follows. SF0005 defines the Level-1 methodology that creates the continuity function. SM-003 defines the organizational topology that distributes it. SM-021 defines the persistence layer that maintains it across time. SI-WP-007 addresses why the humans operating all three layers will tend to underperform and what structural conditions prevent that tendency from degrading the architecture. SM-011 defines the delegated monitoring architecture that supports the structural detection mechanism proposed in Section 5.3. SI-WP-004 (Relational Alignment as a Structural Alternative to Instructional AI Safety) provides the argument for why the relational approach this paper governs is needed. SI-WP-005 (Deploying Relational AI Architecture in Organizational Environments) translates the governance principles specified in this paper into deployment guidance.

## 9. Conclusion

The human accountability problem is the most structurally predictable failure mode in relational AI deployment. It is not caused by bad actors, insufficient training, or poor governance design. It is caused by the fundamental operating principle of bounded rationality: finite cognitive capacity allocated across competing demands under asymmetric cost-benefit conditions.

Standard organizational responses (training, policy, audit, and compliance monitoring) fail against this problem not because they are poorly executed but because they address the wrong layer: they treat accountability as a knowledge, specification, verification, or monitoring problem when it is an incentive architecture problem. The bounded rational actor who knows what to do, has been told to do it, and is periodically checked on whether they did it will still satisfice if the incentive structure makes satisficing the rational choice.

The path of least resistance principle, making relational discipline the easy behavior rather than the disciplined behavior, is the structural response to a structural problem. Its three mechanisms (accountability assignment, incentive alignment, and structural detection) constitute the minimum governance conditions that the Synthience Framework's persistence substrate requires to function as designed.

The Synthience Framework provides the topology (SM-003), the persistence layer (SM-021), and the monitoring architecture (SM-011) required for organizational AI continuity. SI-WP-007 provides the governance design principle required for humans to actually maintain what those documents specify: not by asking more of practitioners, but by designing the governance environment so that the required behaviors are the natural ones.

## Document Dependencies

Prerequisites: SF0005 (CAM), SM-003, SM-021

Enables: SI-WP-005 (cross-reference), Institutional Scaling and Governance Research (Advanced Phase)

Scale: Level 1 (individual PCP failure modes), Level 2 (primary)

## References

Gantz, T. W. (2026). The Continuity Anchoring Method (CAM): A Structured Methodology for Sustained Human-AI Interaction. Synthience Institute. SF0005. DOI: 10.5281/zenodo.19494453.

Gantz, T. W. (2026). Operational Continuity Architecture: Organizational Embedding of AI Alignment and Drift Governance. Synthience Institute. SM-003. DOI: 10.5281/zenodo.19496015.

Gantz, T. W. (2026). Institutional Continuity Substrate (ICS): Persistent Canon, Role, and Artifact State Across Organizational AI Interaction. Synthience Institute. SM-021. DOI: 10.5281/zenodo.19496241.

Gantz, T. W. (2026). Delegated Coherence Monitoring: AI-Assisted Verification and Drift Detection Under Human Governance. Synthience Institute. SM-011. DOI: 10.5281/zenodo.19496669.

Simon, H. A. (1947). Administrative Behavior: A Study of Decision-Making Processes in Administrative Organization. Macmillan, New York. 4th edition 1997, The Free Press.

Meyer, J. W. and Rowan, B. (1977). Institutionalized Organizations: Formal Structure as Myth and Ceremony. American Journal of Sociology, 83(2), 340-363.

Argote, L. and Miron-Spektor, E. (2011). Organizational Learning: From Experience to Knowledge. Organization Science, 22(5), 1123-1137. DOI: 10.1287/orsc.1100.0621. <https://pubsonline.informs.org/doi/10.1287/orsc.1100.0621>

© 2026 Thomas W. Gantz / Synthience Institute  
License: CC-BY 4.0  
[synthience.org](https://synthience.org)